



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in **Computer Architecture** (a.k.a. Machine Structures)

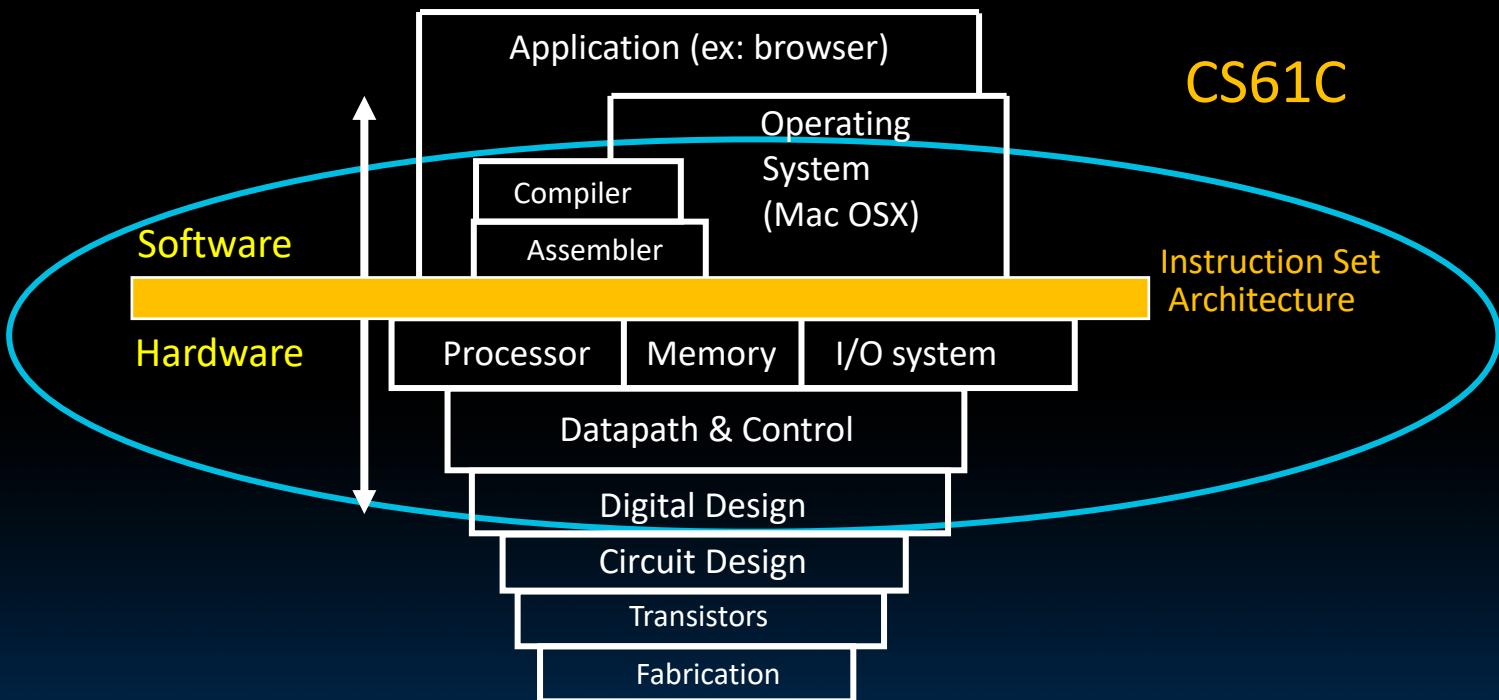


UC Berkeley
Professor
Bora Nikolić

RISC-V Processor Design

Machine Structures

CS61C



New-School Machine Structures

Software

Parallel Requests

Assigned to computer

e.g., Search "Cats"

Parallel Threads

Assigned to core e.g., Lookup, Ads

Parallel Instructions

>1 instruction @ one time

e.g., 5 pipelined instructions

Parallel Data

>1 data item @ one time

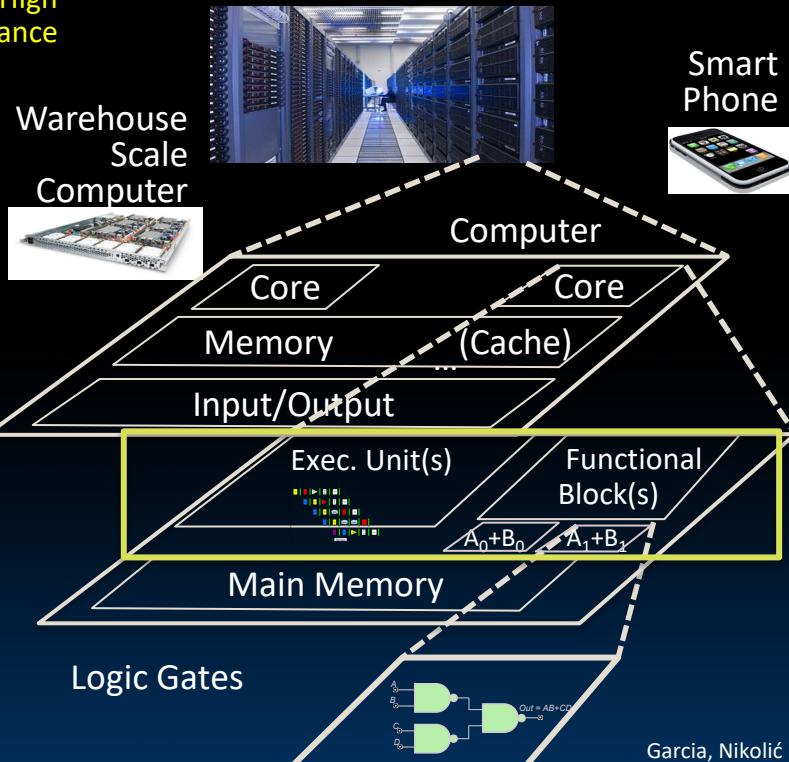
e.g., Add of 4 pairs of words

Hardware descriptions

All gates work in parallel at same time

Harness
Parallelism &
Achieve High
Performance

Hardware



Great Idea #1: Abstraction (Levels of Representation/Interpretation)

High Level Language
Program (e.g., C)

| Compiler

Assembly Language
Program (e.g., RISC-V)

| Assembler

Machine Language
Program (RISC-V)

```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

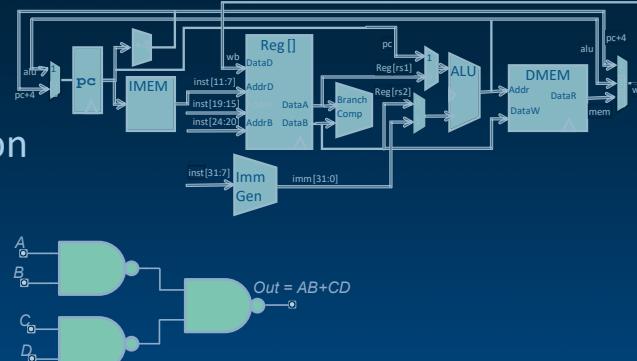
```
lw    x3, 0(x10)  
lw    x4, 4(x10)  
sw    x4, 0(x10)  
sw    x3, 4(x10)
```

1000	1101	1110	0010	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	
1000	1110	0001	0000	0000	0000	0000	0000	0000	0000	0000	0000	0100				
1010	1110	0001	0010	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	
1010	1101	1110	0010	0000	0000	0000	0000	0000	0000	0000	0000	0100				

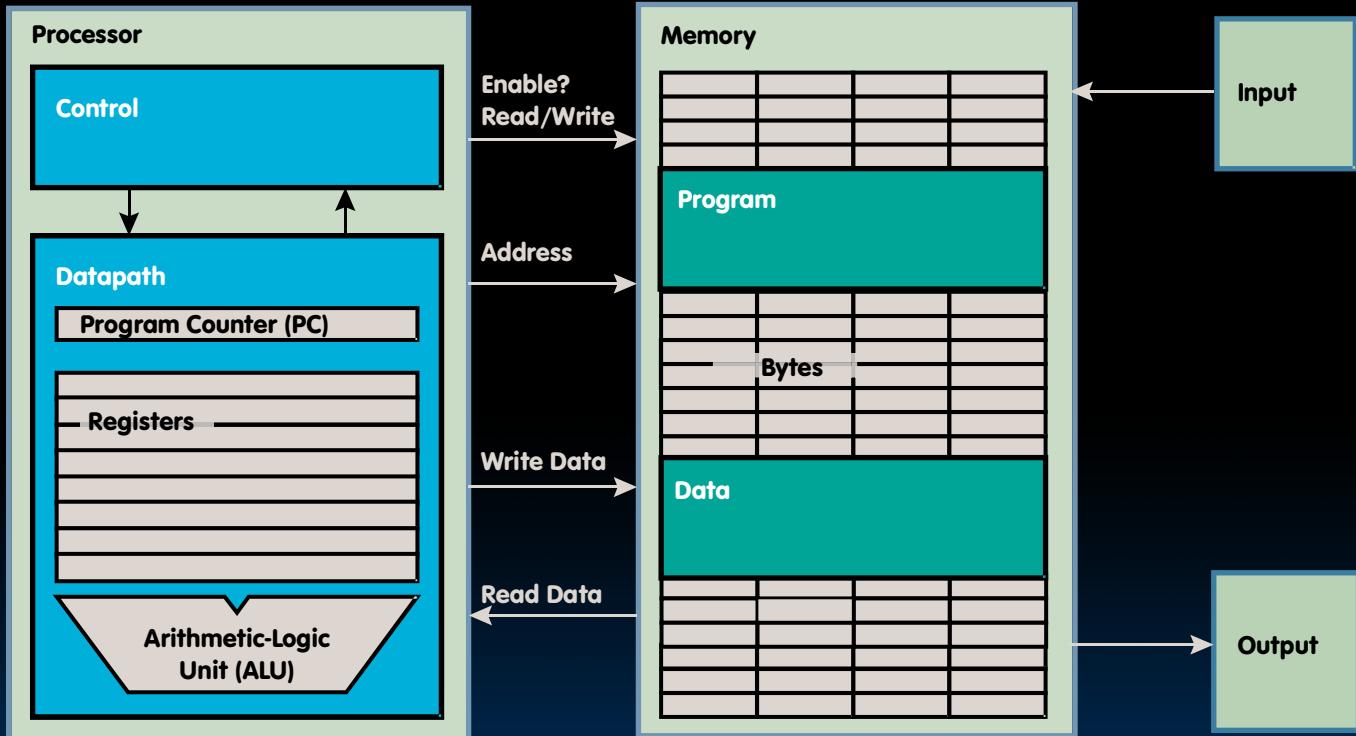
Hardware Architecture Description
(e.g., block diagrams)

| Architecture Implementation

Logic Circuit Description
(Circuit Schematic Diagrams)



Our Single-Core Processor So Far...



- Processor (CPU): the active part of the computer that does all the work (data manipulation and decision-making)
- Datapath: portion of the processor that contains hardware necessary to perform operations required by the processor (the brawn)
- Control: portion of the processor (also in hardware) that tells the datapath what needs to be done (the brain)

Need to Implement All RV32I Instructions

OpenRISC RISC-V Reference Card

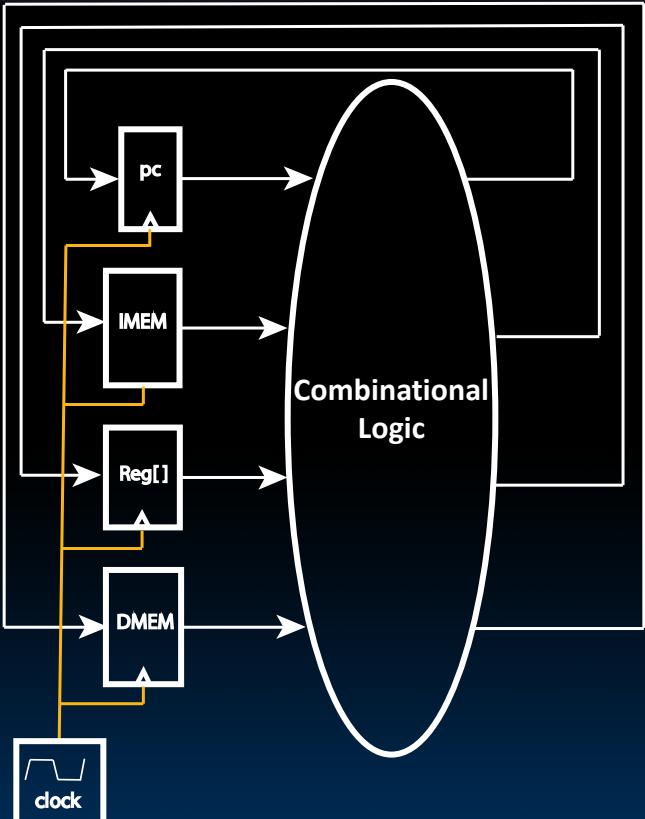
Base Integer Instructions: RV32I									
Category	Name	Fmt	RV32I Base	Category	Name	Fmt	RV32I Base		
Shifts	Shift Left Logical	R	SLL rd,rs1,rs2	Loads	Load Byte	I	LB rd,rs1,imm		
	Shift Left Log. Imm.	I	SLLI rd,rs1,shamt		Load Halfword	I	LH rd,rs1,imm		
	Shift Right Logical	R	SRL rd,rs1,rs2		Load Byte Unsigned	I	LBU rd,rs1,imm		
	Shift Right Log. Imm.	I	SRLI rd,rs1,shamt		Load Half Unsigned	I	LHU rd,rs1,imm		
	Shift Right Arithmetic	R	SRA rd,rs1,rs2		Load Word	I	LW rd,rs1,imm		
	Shift Right Arith. Imm.	I	SRAI rd,rs1,shamt	Stores	Store Byte	S	SB rs1,rs2,imm		
Arithmetic	ADD	R	ADD rd,rs1,rs2		Store Halfword	S	SH rs1,rs2,imm		
	ADD Immediate	I	ADDI rd,rs1,imm		Store Word	S	SW rs1,rs2,imm		
	SUBtract	R	SUB rd,rs1,rs2	Branches	Branch =	B	BEQ rs1,rs2,imm		
	Load Upper Imm	U	LUI rd,imm		Branch ≠	B	BNE rs1,rs2,imm		
	Add Upper Imm to PC	U	AUIPC rd,imm		Branch <	B	BLT rs1,rs2,imm		
Logical	XOR	R	XOR rd,rs1,rs2		Branch ≥	B	BGE rs1,rs2,imm		
	XOR Immediate	I	XORI rd,rs1,imm		Branch < Unsigned	B	BLTU rs1,rs2,imm		
	OR	R	OR rd,rs1,rs2		Branch ≥ Unsigned	B	BGEU rs1,rs2,imm		
	OR Immediate	I	ORI rd,rs1,imm	Jump & Link	J&L	J	JAL rd,imm		
	AND	R	AND rd,rs1,rs2		Jump & Link Register	I	JALR rd,rs1,imm		
	AND Immediate	I	ANDI rd,rs1,imm						
Compare	Set <	R	SLT rd,rs1,rs2	Synch	Synch thread	I	FENCE		
	Set < Immediate	I	SLTI rd,rs1,imm						
	Set < Unsigned	R	SLTU rd,rs1,rs2	Environment	CALL	I	ECALL		
	Set < Imm Unsigned	I	SLTIU rd,rs1,imm		BREAK	I	EBREAK		



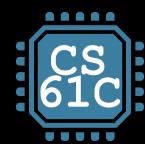
Not in
61C

Building a RISC-V Processor

One-Instruction-Per-Cycle RISC-V Machine



- On every tick of the clock, the computer executes one instruction
- Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
- At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle



Stages of the Datapath : Overview

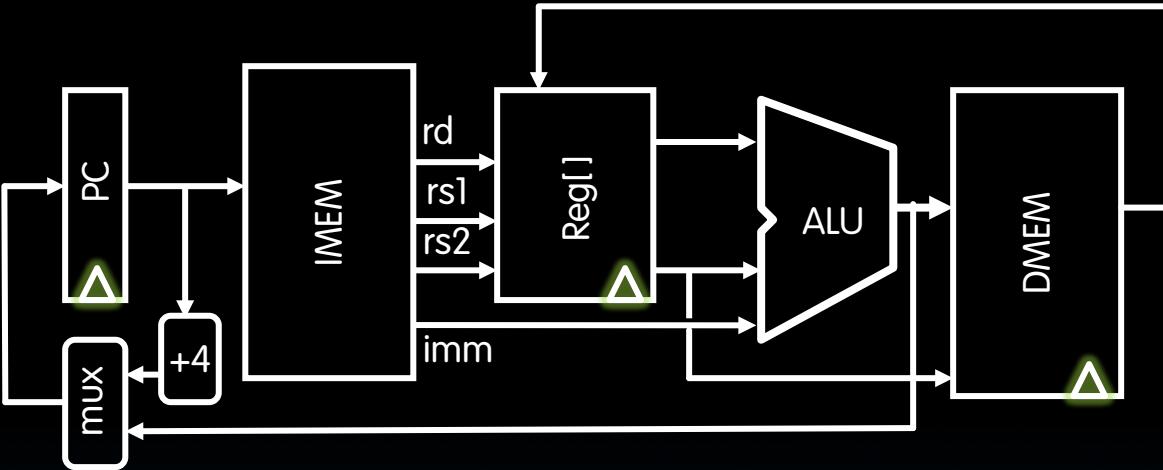
- Problem: a single, “monolithic” block that “executes an instruction” (performs all necessary operations beginning with fetching the instruction) would be too bulky and inefficient
- Solution: break up the process of “executing an instruction” into stages, and then connect the stages to create the whole datapath
 - smaller stages are easier to design
 - easy to optimize (change) one stage without touching the others (modularity)



Five Stages of the Datapath

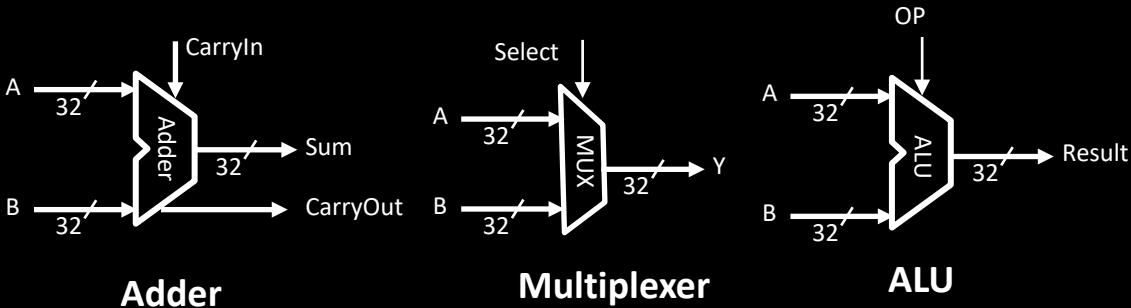
- Stage 1: *Instruction Fetch (IF)*
- Stage 2: *Instruction Decode (ID)*
- Stage 3: *Execute (EX) - ALU (Arithmetic-Logic Unit)*
- Stage 4: *Memory Access (MEM)*
- Stage 5: *Write Back to Register (WB)*

Basic Phases of Instruction Execution



Datapath Components: Combinational

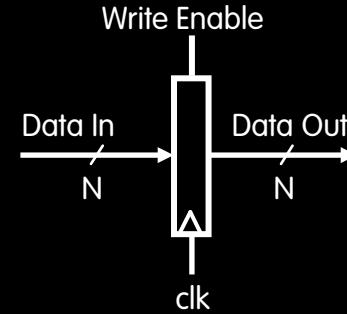
- Combinational elements



- Storage elements + clocking methodology
- Building blocks

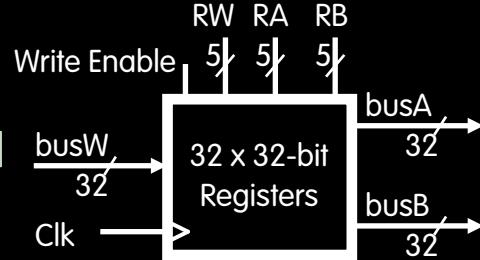
Datapath Elements: State and Sequencing (1/3)

- Register
- Write Enable:
 - Low (or deasserted) (0): Data Out will not change
 - Asserted (1): Data Out will become Data In on positive edge of clock



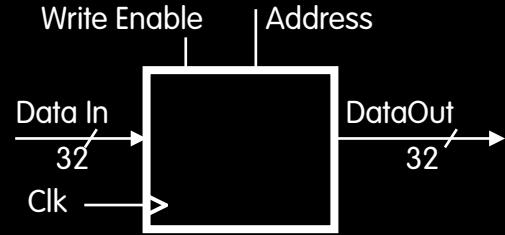
Datapath Elements: State and Sequencing (2/3)

- Register file (regfile, RF) consists of 32 registers:
 - Two 32-bit output busses: busA and busB
 - One 32-bit input bus: busW
- Register is selected by:
 - RA (number) selects the register to put on busA (data)
 - RB (number) selects the register to put on busB (data)
 - RW (number) selects the register to be written via busW (data) when Write Enable is 1
- Clock input (Clk)
 - Clk input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block:
 - RA or RB valid \Rightarrow busA or busB valid after “access time.”



Datapath Elements: State and Sequencing (3/3)

- “Magic” Memory
 - One input bus: Data In
 - One output bus: Data Out
- Memory word is found by:
 - For Read: Address selects the word to put on Data Out
 - For Write: Set Write Enable = 1: address selects the memory word to be written via the Data In bus
- Clock input (CLK)
 - CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block: Address valid \Rightarrow Data Out valid after “access time”



State Required by RV32I ISA (1/2)

Each instruction during execution reads and updates the state of : (1) Registers, (2) Program counter, (3) Memory

- Registers (**x0 . . x31**)
 - Register file (*regfile*) **Reg** holds 32 registers x 32 bits/register:
Reg [0] . . Reg [31]
 - First register read specified by **rs1** field in instruction
 - Second register read specified by **rs2** field in instruction
 - Write register (destination) specified by *rd* field in instruction
 - **x0** is always 0 (writes to **Reg [0]** are ignored)
- Program Counter (**PC**)
 - Holds address of current instruction



State Required by RV32I ISA (2/2)

- Memory (**MEM**)
 - Holds both instructions & data, in one 32-bit byte-addressed memory space
 - We'll use separate memories for instructions (**IMEM**) and data (**DMEM**)
 - *These are placeholders for instruction and data caches*
 - Instructions are read (*fetched*) from instruction memory (assume **IMEM** read-only)
 - Load/store instructions access data memory

**R-Type Add
Datapath**

Review: R-Type Instructions

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	10	9	8	7	6	5	4	3	2	1	0
R-format : ALU																														
[31:25]					[24:20]					[19:15]					[14:12]				[11:7]				[6:0]							
7					5					5					3				5				7							
func7					rs2					rs1					func3				rd				opcode							
0000000					rs2			rs1				000	: ADD	rd										0110011:OP-R						
0100000					rs2			rs1				000	: SUB	rd										0110011:OP-R						
0000000					rs2			rs1				001	: SLL	rd										0110011:OP-R						
0000000					rs2			rs1				010	: SLT	rd										0110011:OP-R						
0000000					rs2			rs1				011	: SLTU	rd										0110011:OP-R						
0000000					rs2			rs1				100	: XOR	rd										0110011:OP-R						
0000000					rs2			rs1				101	: SRL	rd										0110011:OP-R						
0100000					rs2			rs1				101	: SRA	rd										0110011:OP-R						
0000000					rs2			rs1				110	: OR	rd										0110011:OP-R						
0000000					rs2			rs1				111	: AND	rd										0110011:OP-R						

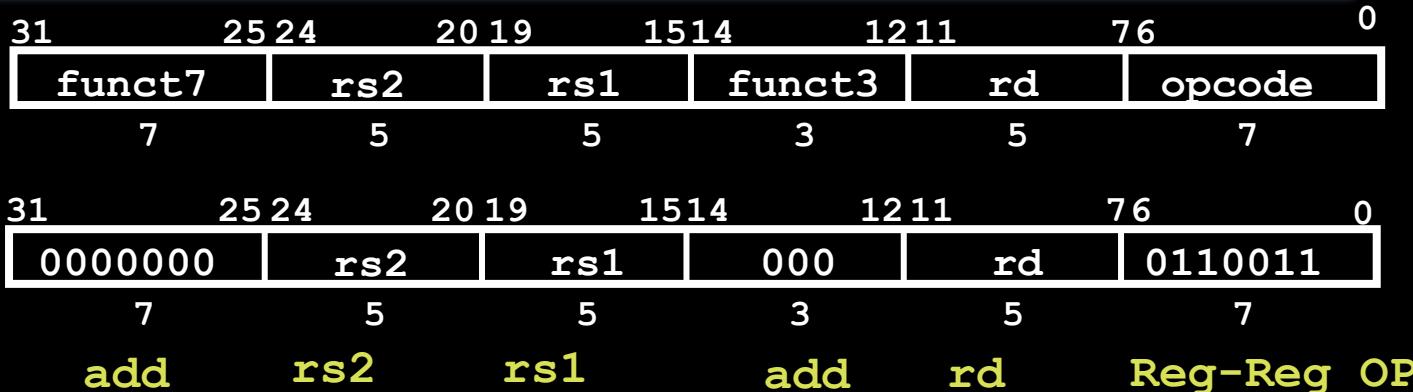
- E.g. Addition/subtraction **add rd, rs1, rs2**

$$R[rd] = R[rs1] + R[rs2]$$

sub rd, rs1, rs2

$$R[rd] = R[rs1] - R[rs2]$$

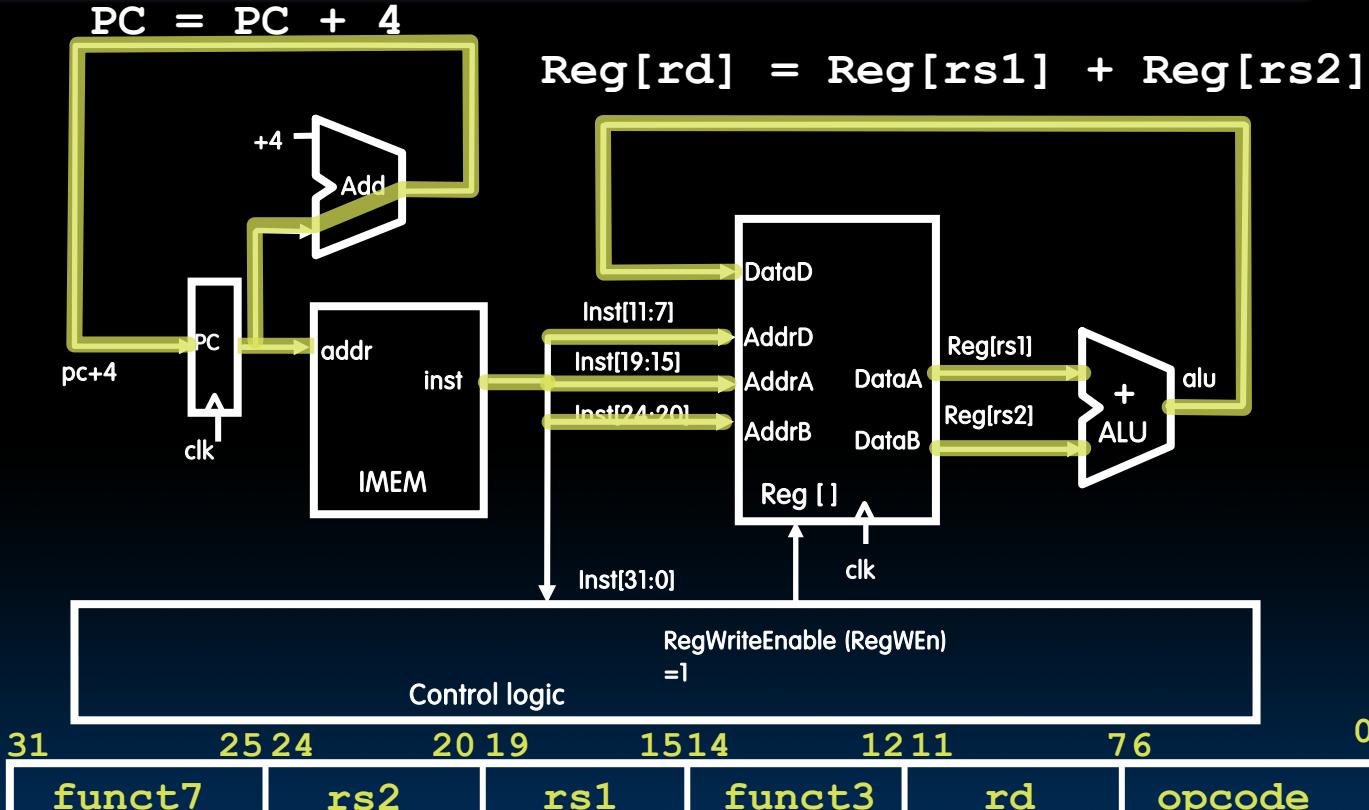
Implementing the add instruction



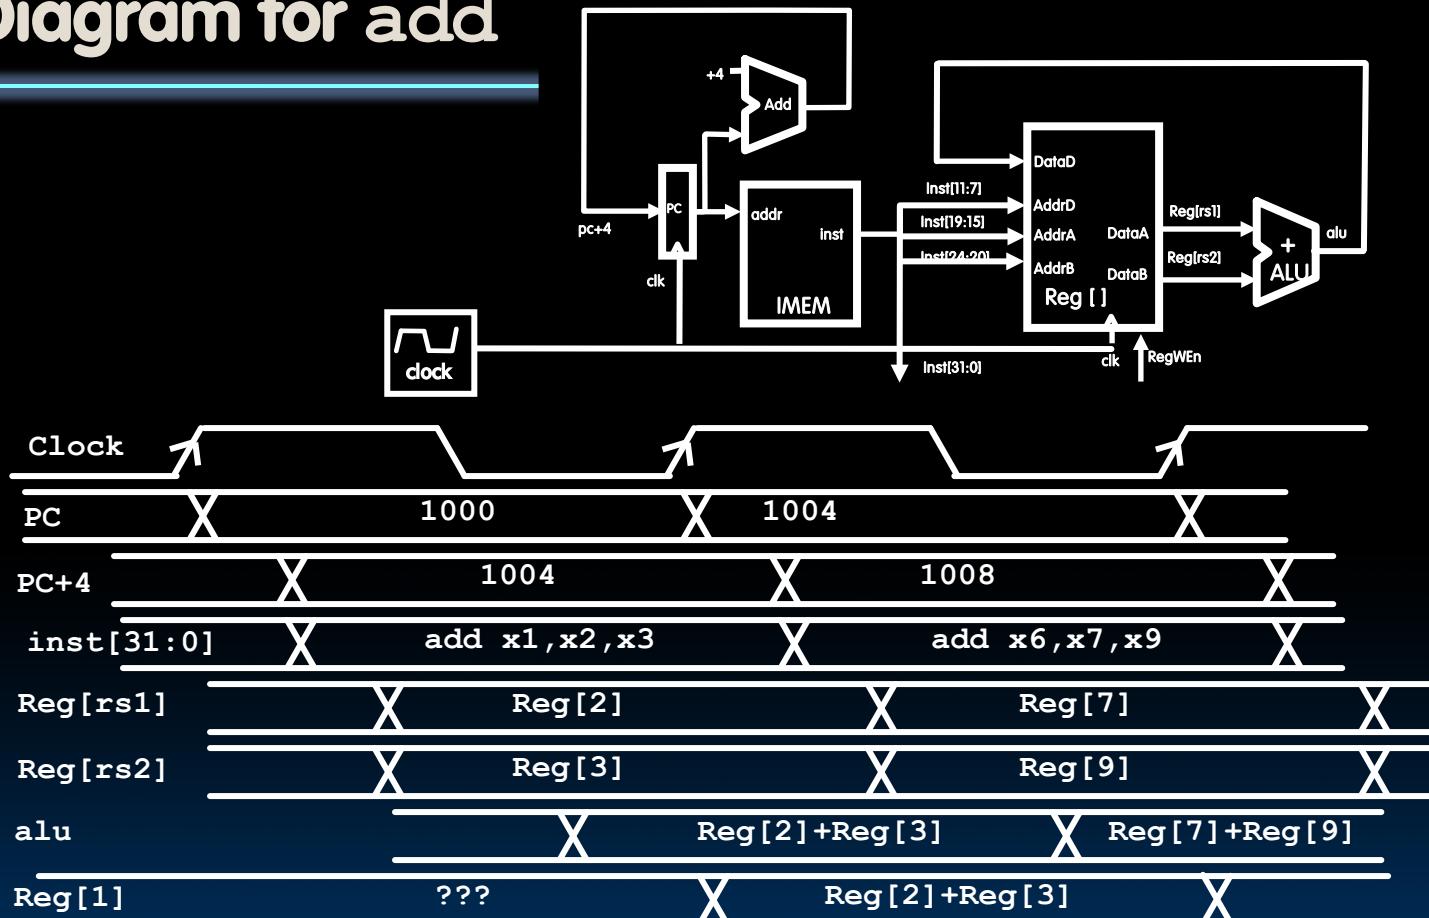
add rd, rs1, rs2

- Instruction makes two changes to machine's state:
 - $\text{Reg}[rd] = \text{Reg}[rs1] + \text{Reg}[rs2]$
 - $\text{PC} = \text{PC} + 4$

Datapath for add



Timing Diagram for add



Sub Datapath

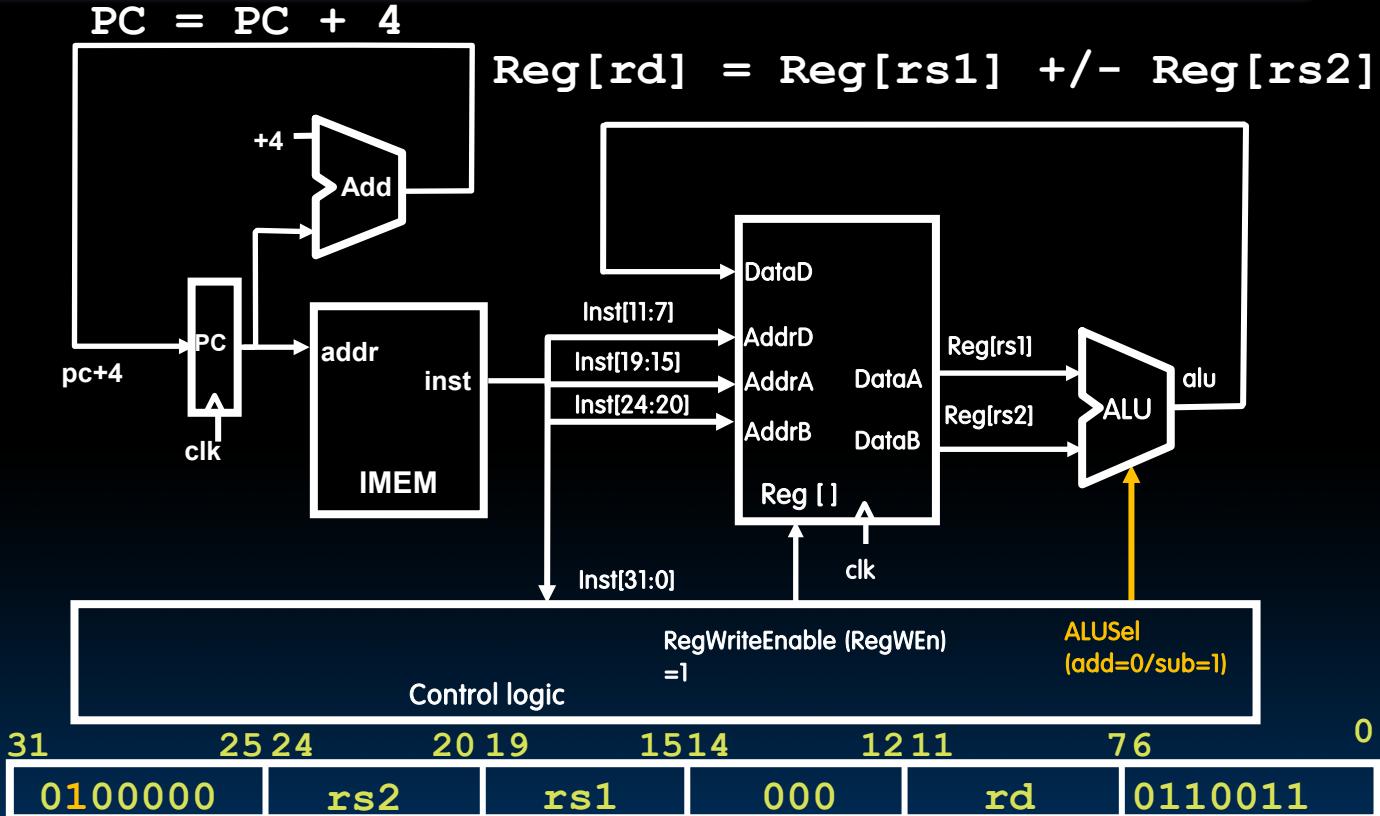
Implementing the sub instruction

0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub

sub rd, rs1, rs2

- Almost the same as add, except now have to subtract operands instead of adding them
- **inst[30]** selects between add and subtract

Datapath for add/sub



Implementing Other R-Format Instructions

0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	001	rd	0110011	sll
0000000	rs2	rs1	010	rd	0110011	slt
0000000	rs2	rs1	011	rd	0110011	sltu
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	101	rd	0110011	srl
0100000	rs2	rs1	101	rd	0110011	sra
0000000	rs2	rs1	110	rd	0110011	or
0000000	rs2	rs1	111	rd	0110011	and

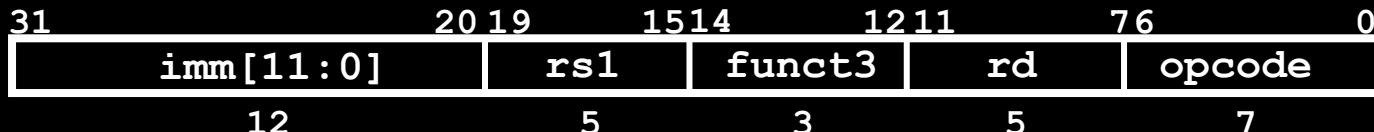
All implemented by decoding funct3 and funct7 fields
and selecting appropriate ALU function

Datapath With Immediates

Implementing I-Format - addi instruction

- RISC-V Assembly Instruction:

addi x15,x1,-50



imm=-50

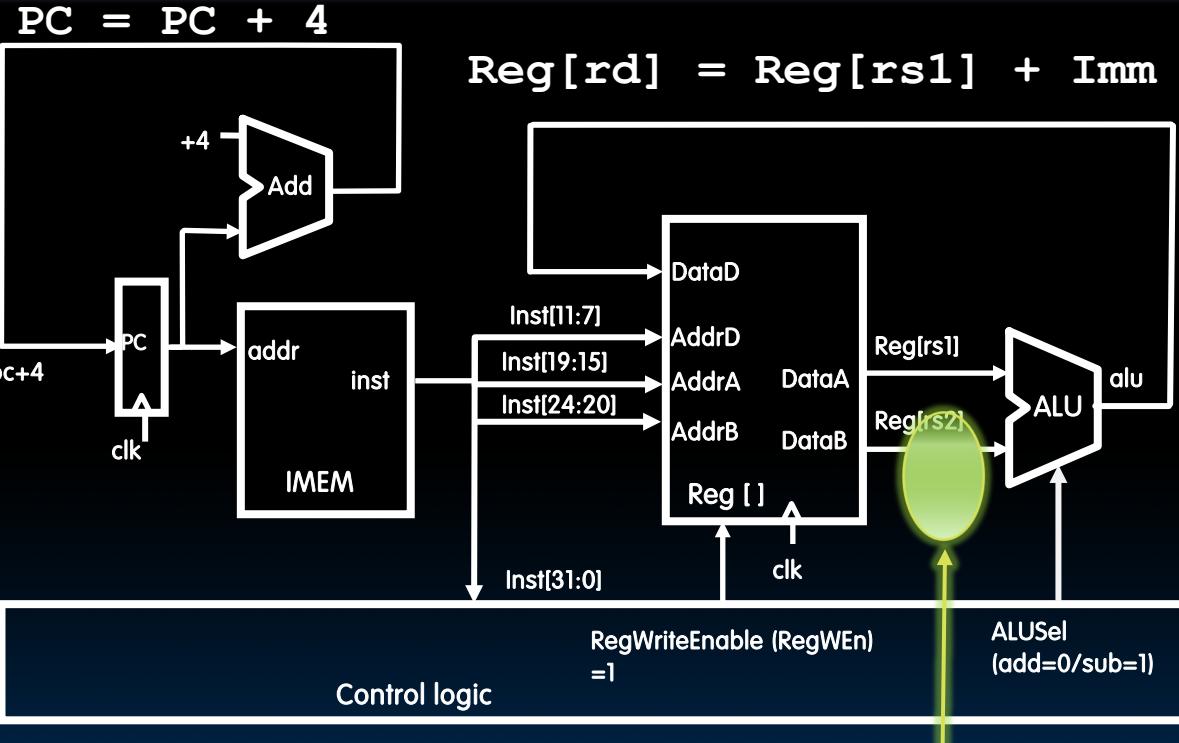
rs1=1

add

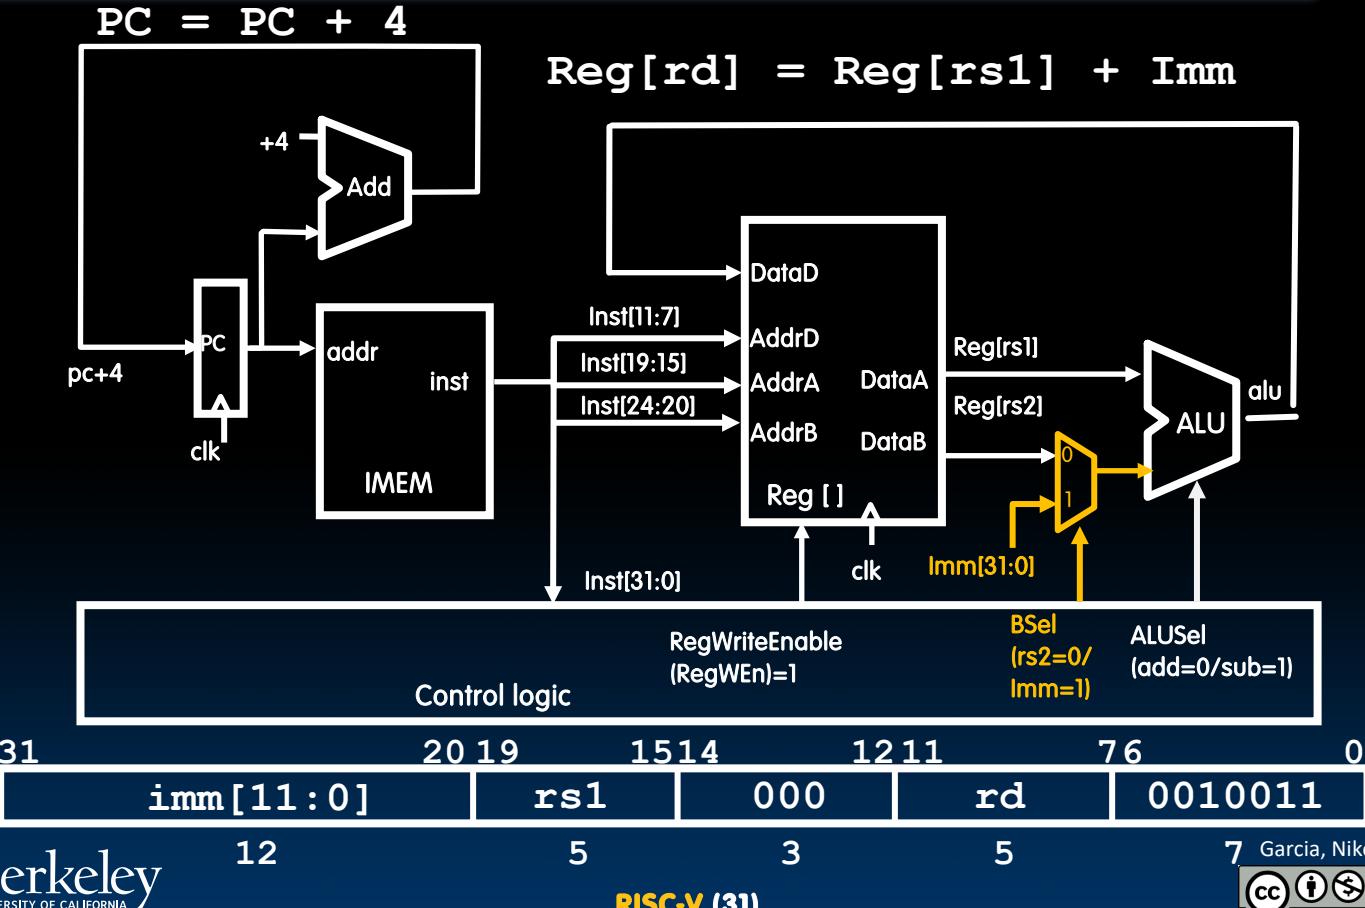
rd=15

OP-Imm

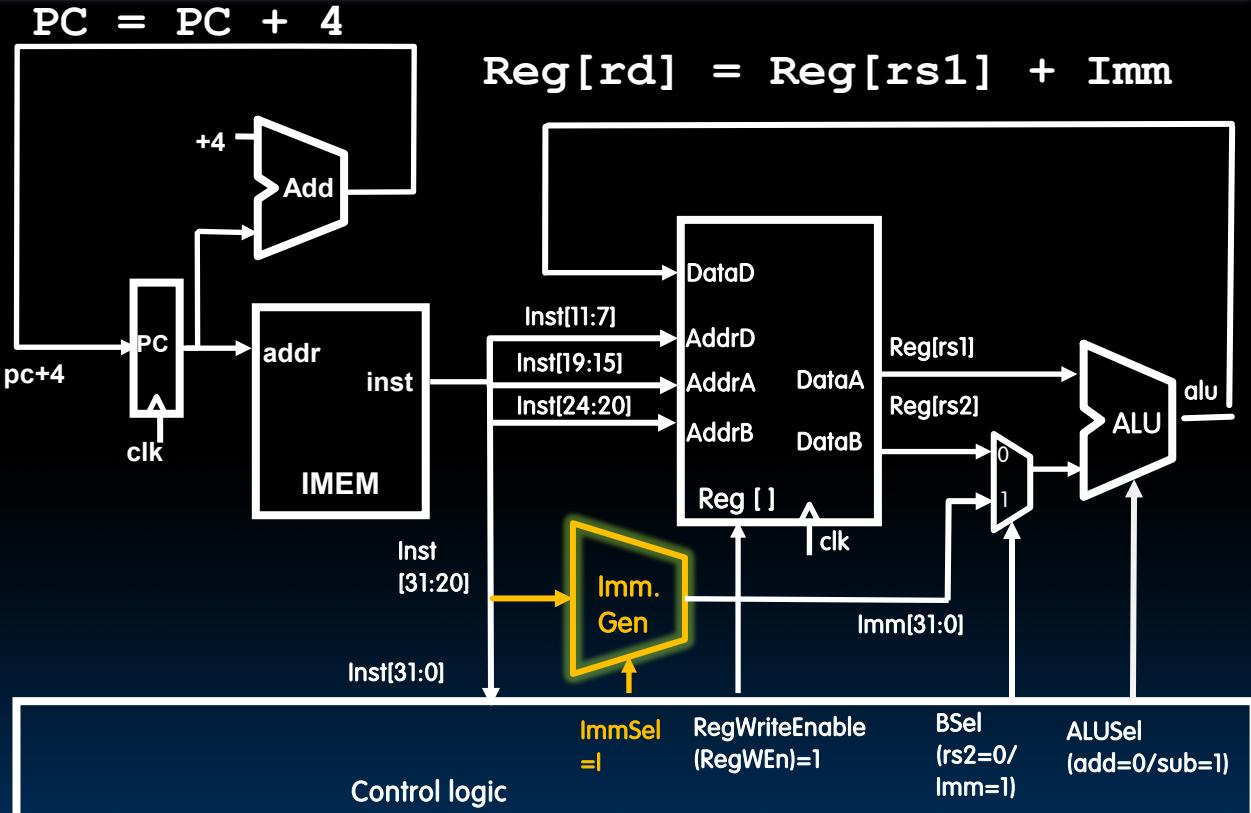
Datapath for add/sub



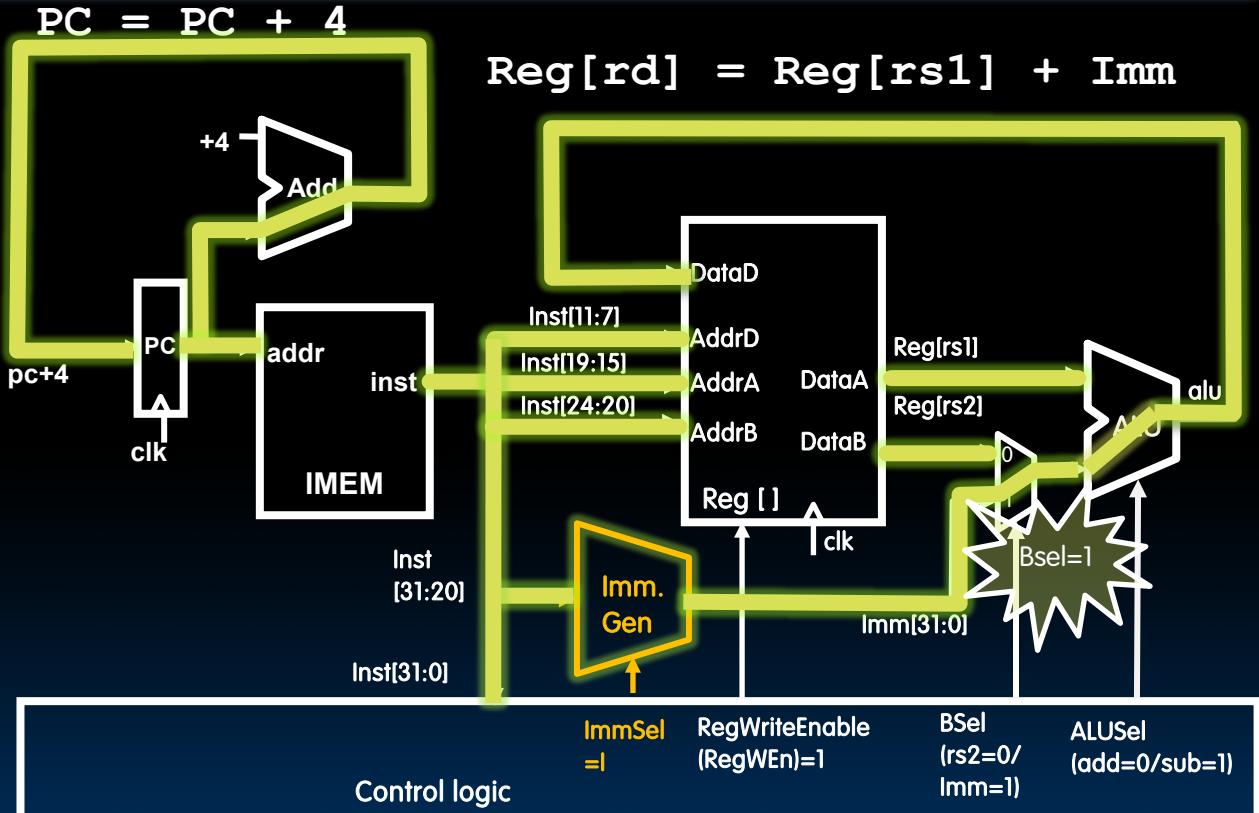
Adding addi to Datapath



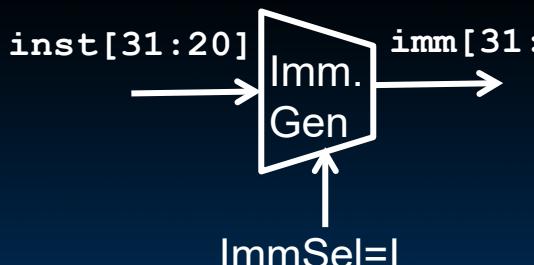
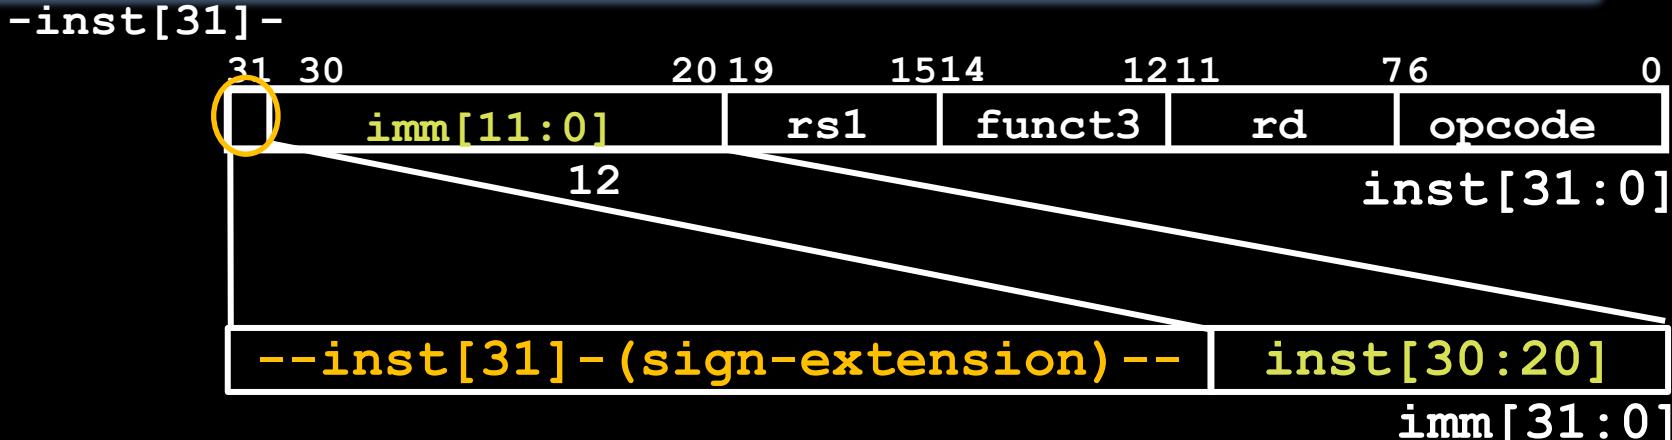
Adding addi to Datapath



Adding addi to Datapath

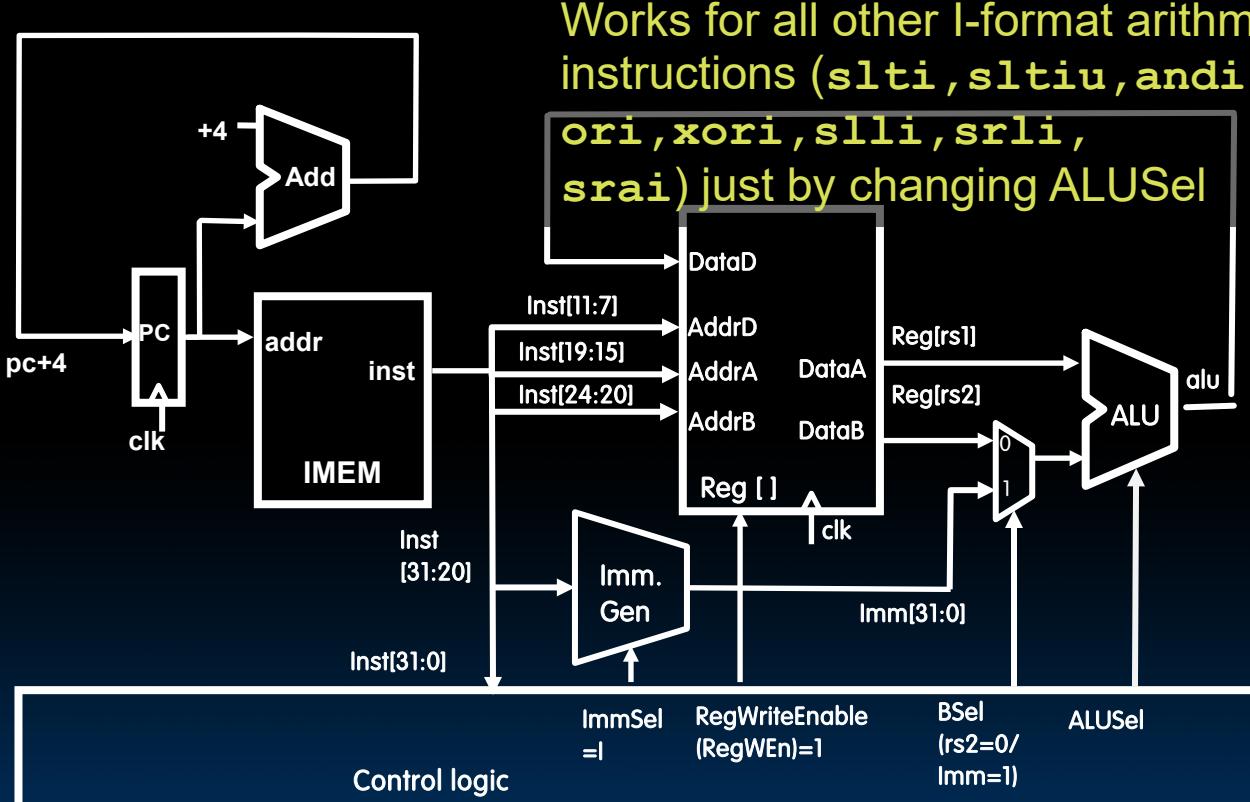


I-Format Immediates



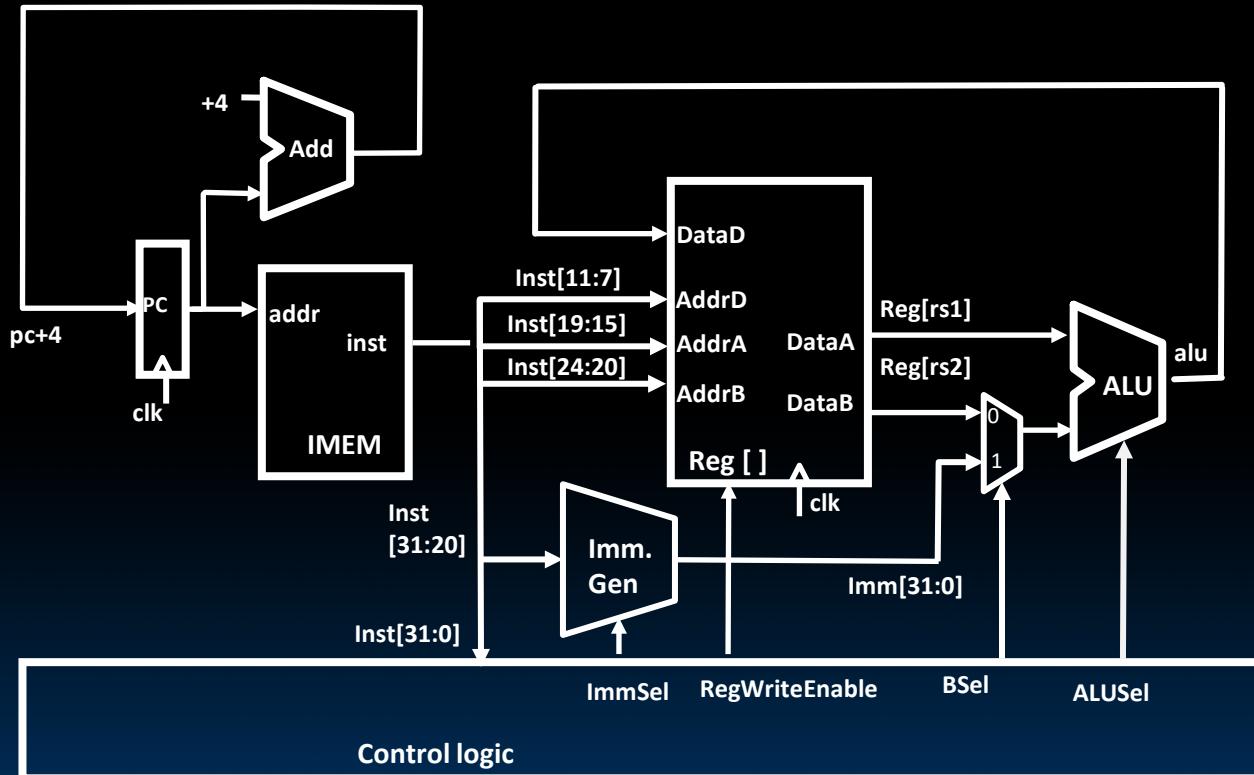
- High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])
- Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])

Adding addi to Datapath



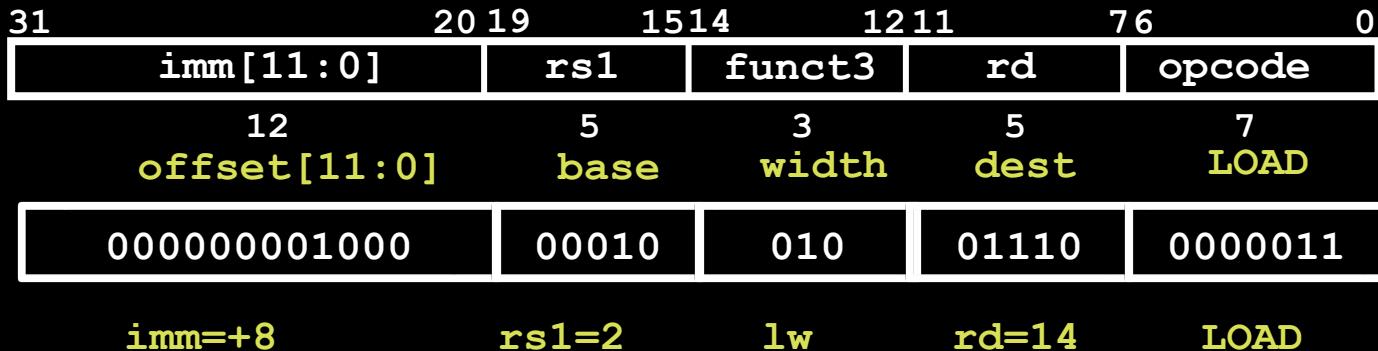
Supporting Loads

R+I Arithmetic/Logic Datapath



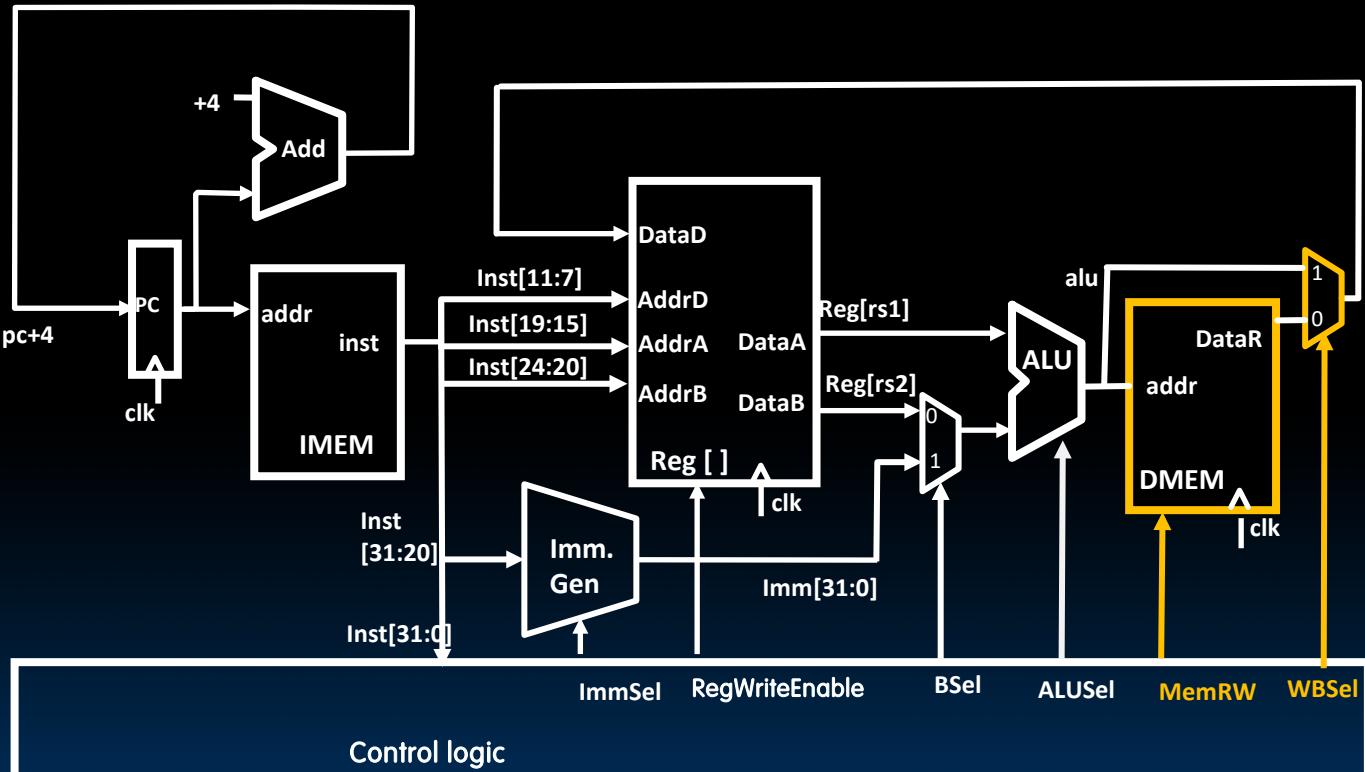
Add lw

- RISC-V Assembly Instruction (I-type): **lw x14, 8(x2)**

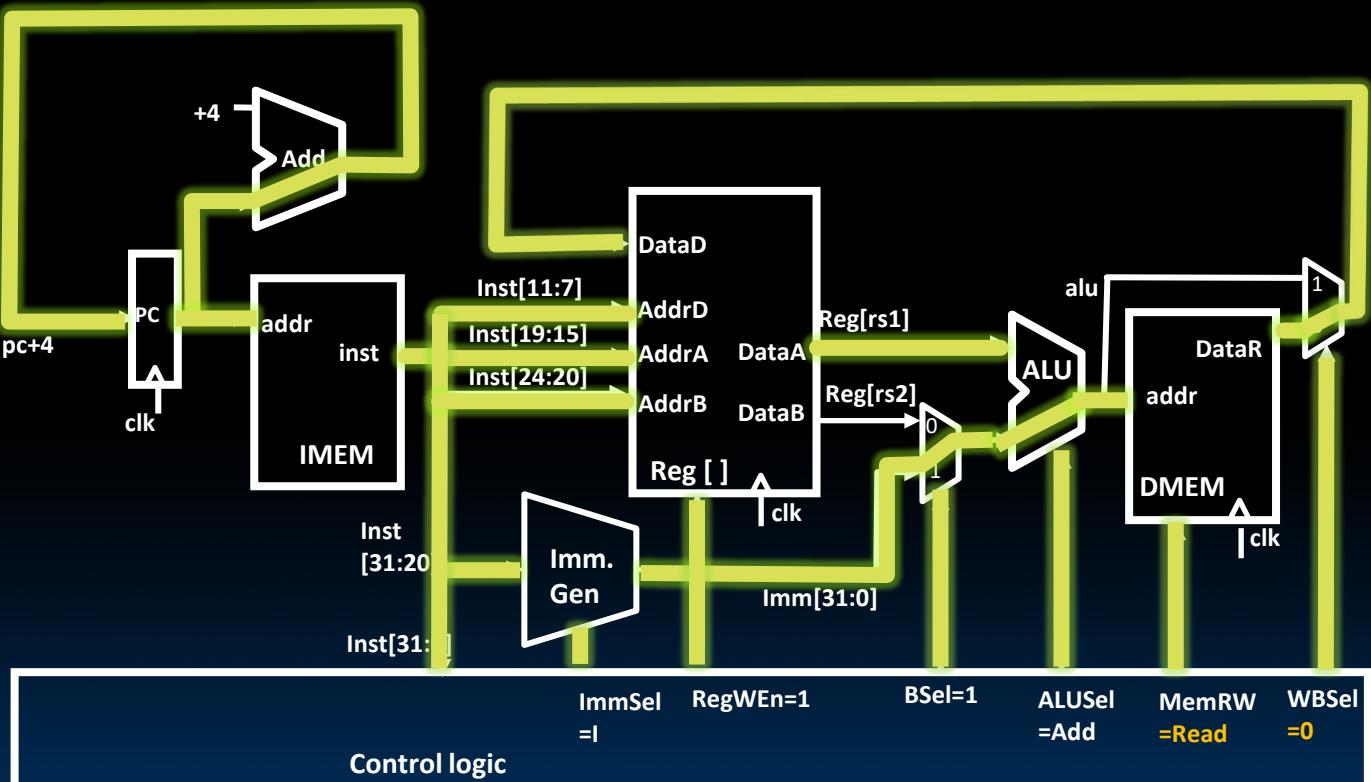


- The 12-bit signed immediate is added to the base address in register **rs1** to form the **memory** address
 - This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from **memory** is stored in register **rd**

R+I Arithmetic/Logic Datapath



R+I Arithmetic/Logic Datapath



All RV32 Load Instructions

imm[11:0]	rs1	000	rd	0000011	lb
imm[11:0]	rs1	001	rd	0000011	lh
imm[11:0]	rs1	010	rd	0000011	lw
imm[11:0]	rs1	100	rd	0000011	lbu
imm[11:0]	rs1	101	rd	0000011	lhu

funct3 field encodes size and 'signedness' of load data

- Supporting the narrower loads requires additional logic to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to register file.
 - It is just a mux + a few gates

Datapath for Stores

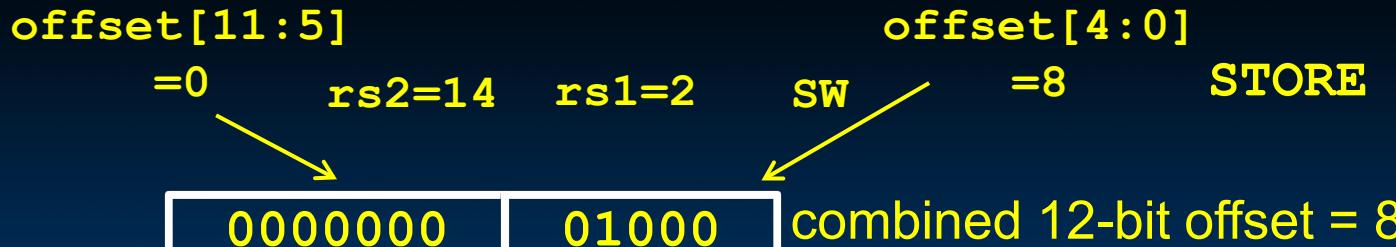
Adding sw instruction

- **sw**: Reads two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!

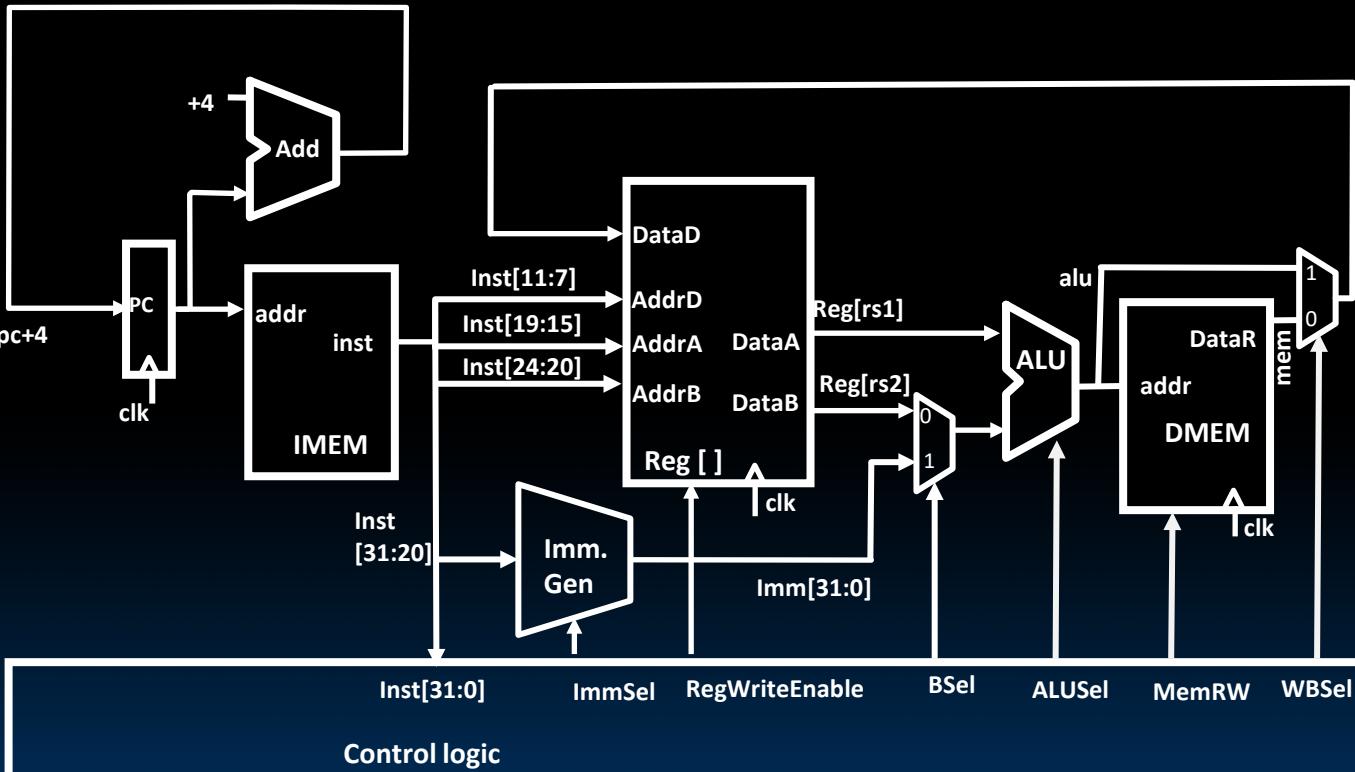
sw x14, 8(x2)

31	25 24	20 19	15 14	12 11	7 6	0
Imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode	
offset[11:5]	src	base	width	offset[4:0]	STORE	

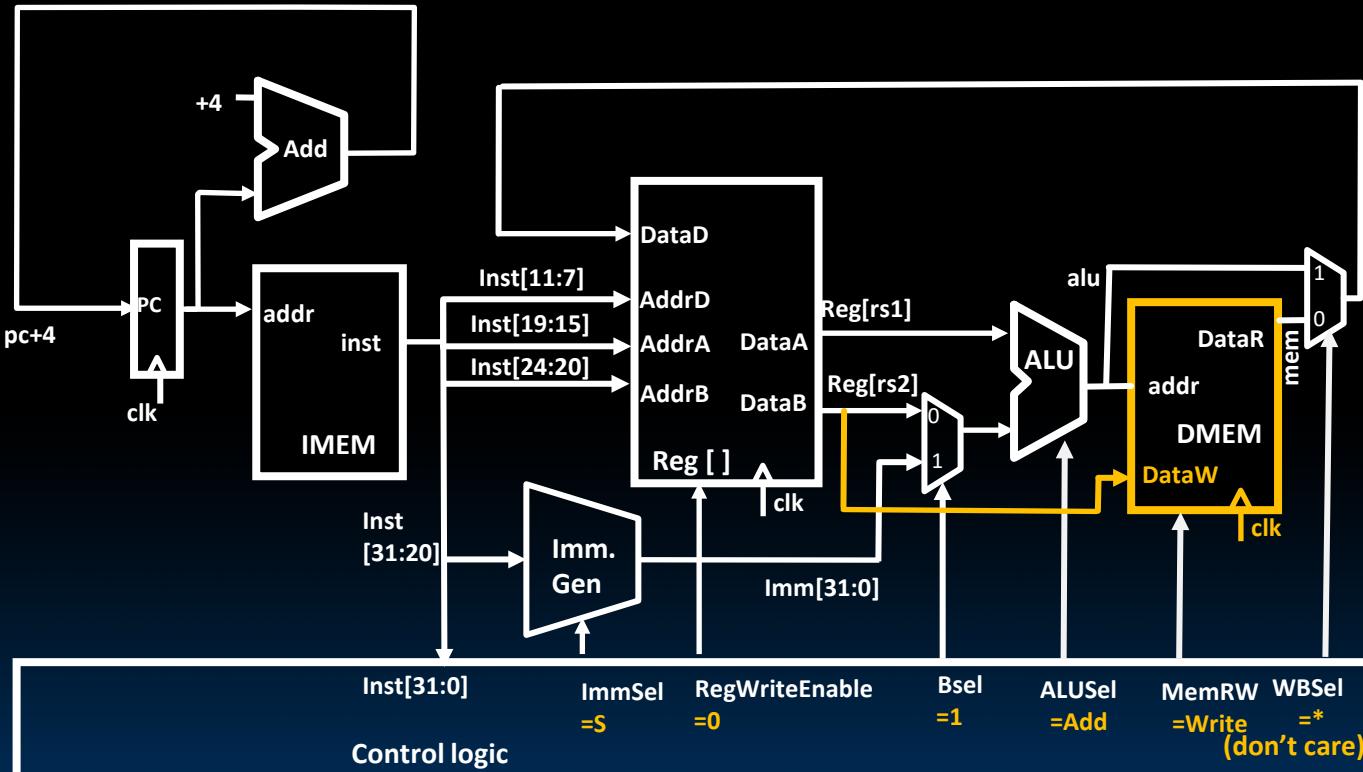
0000000	01110	00010	010	01000	0100011
---------	-------	-------	-----	-------	---------



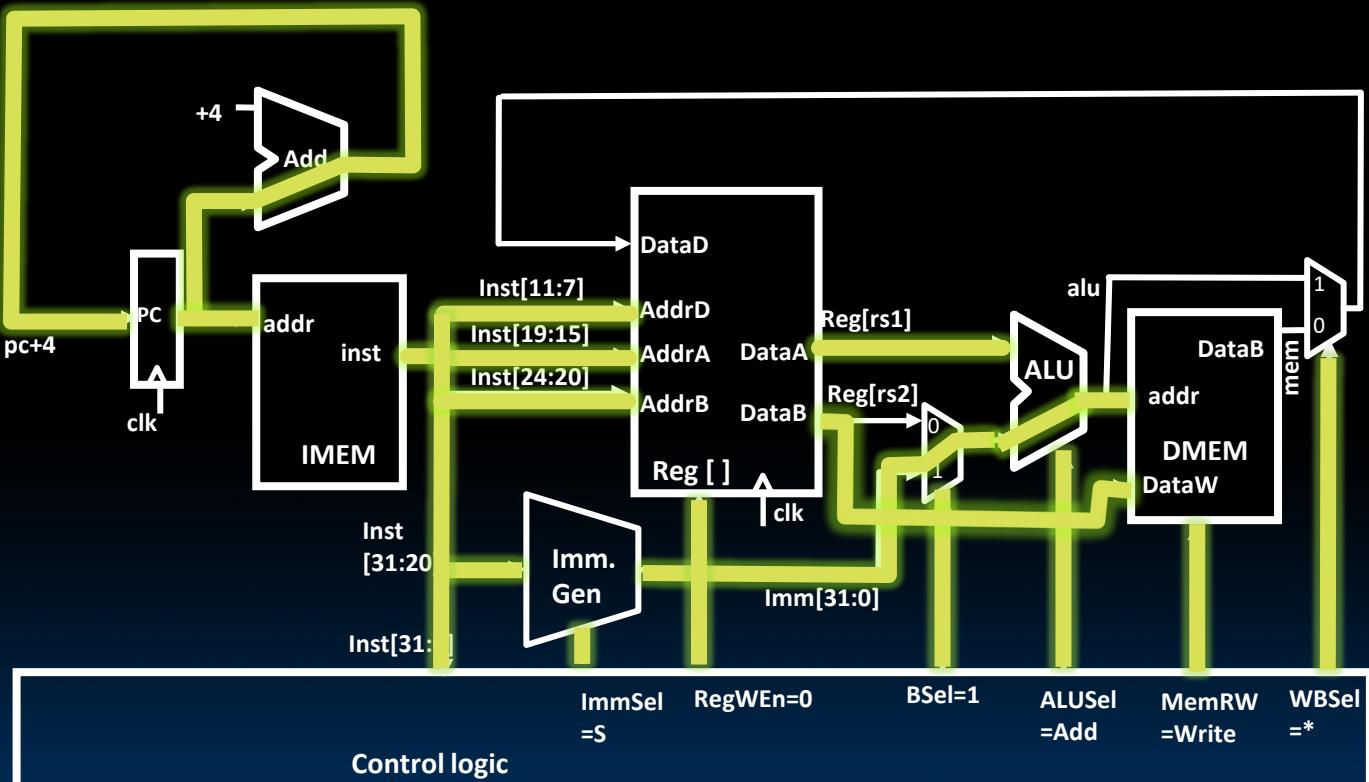
Datapath with 1w



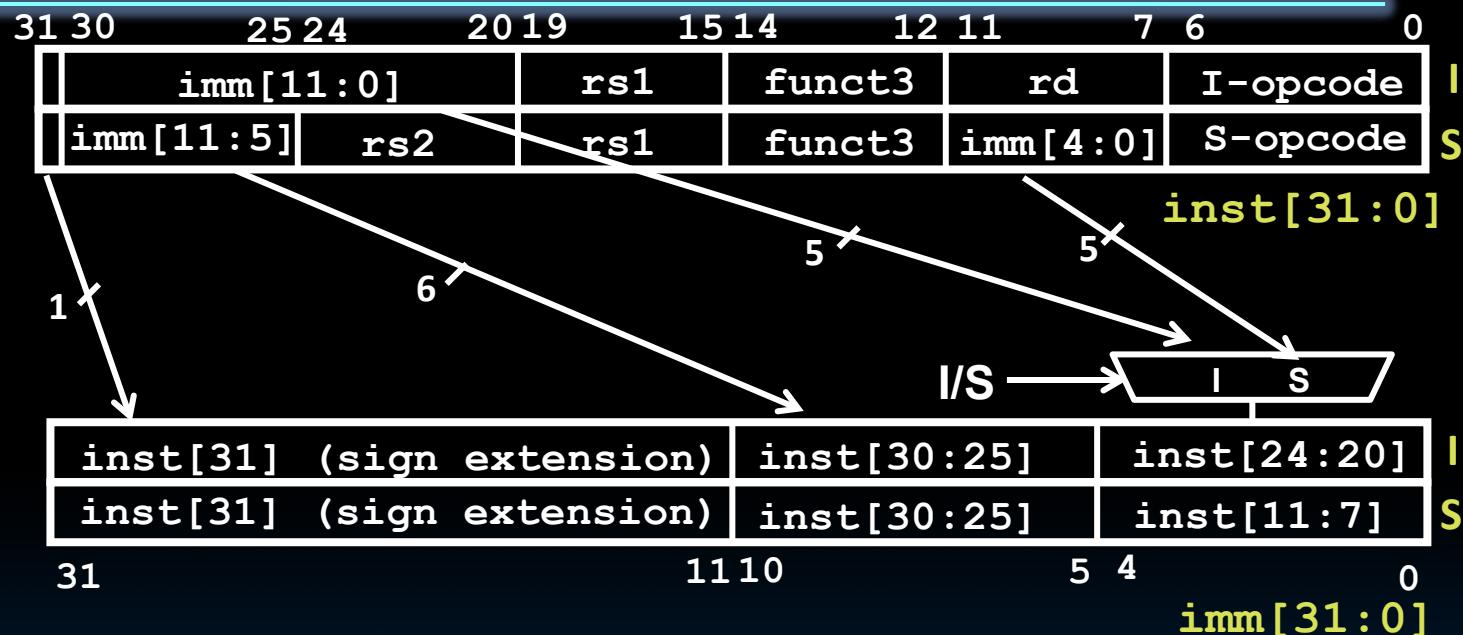
Adding sw to Datapath



Adding sw to Datapath



I+S Immediate Generation



- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
 - Other bits in immediate are wired to fixed positions in instruction

All RV32 Store Instructions

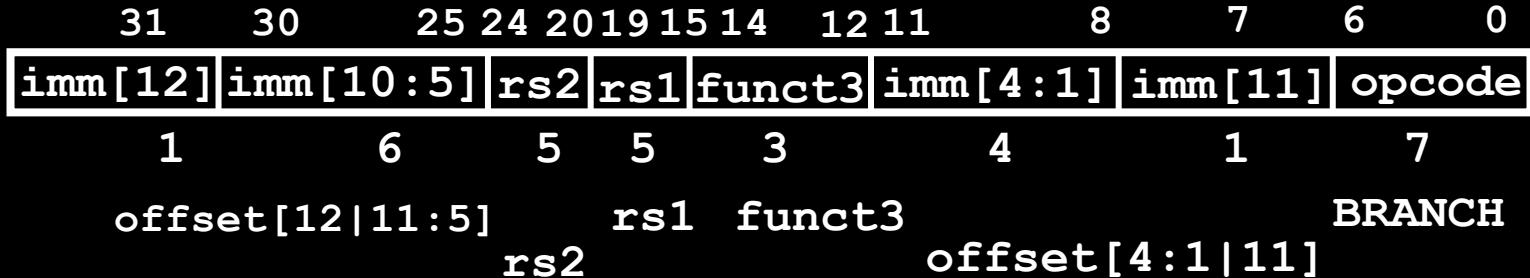
- Store byte writes the low byte to memory
- Store halfword writes the lower two bytes to memory

Imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	sb
Imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	sh
Imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw

width

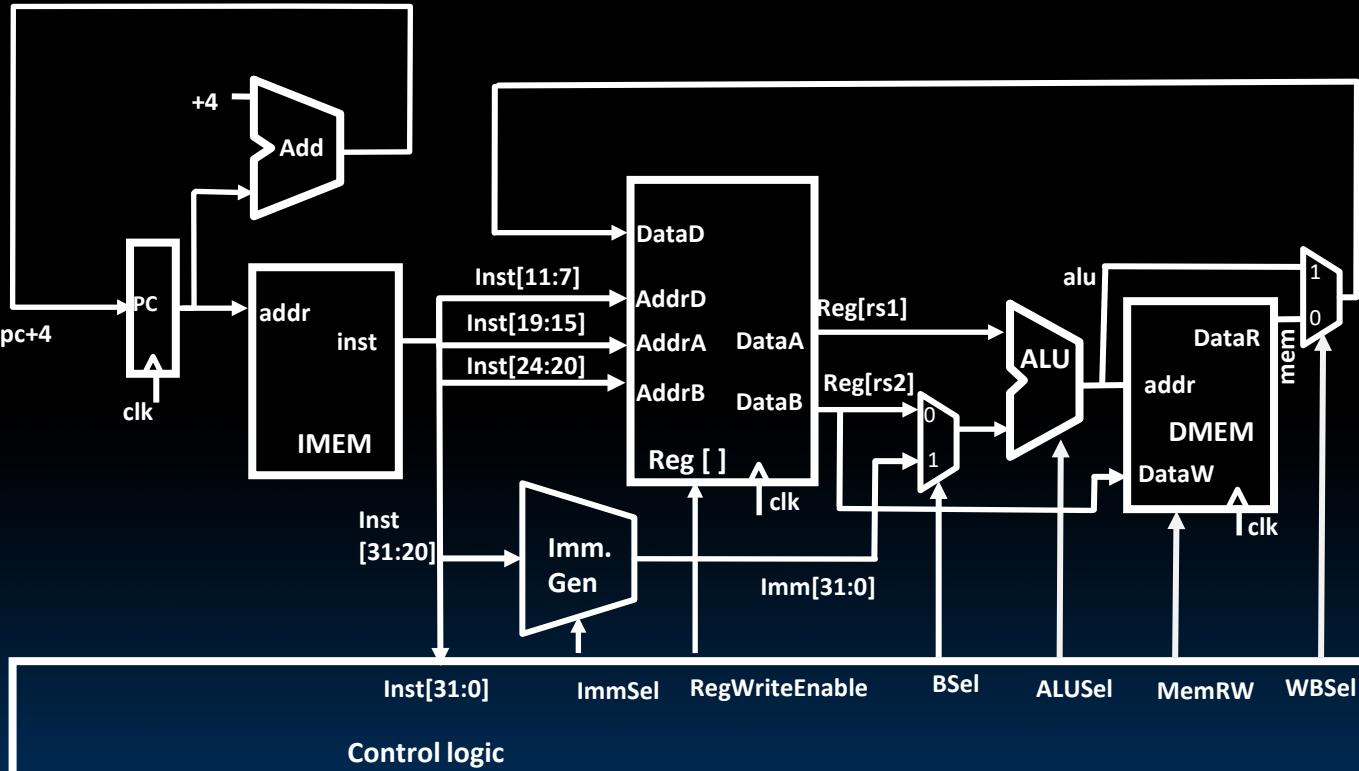
Implementing Branches

RISC-V B-Format for Branches



- B-format is mostly same as S-Format, with two register sources (**rs1/rs2**) and a 12-bit immediate **imm[12:1]**
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode even 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

Datapath So Far



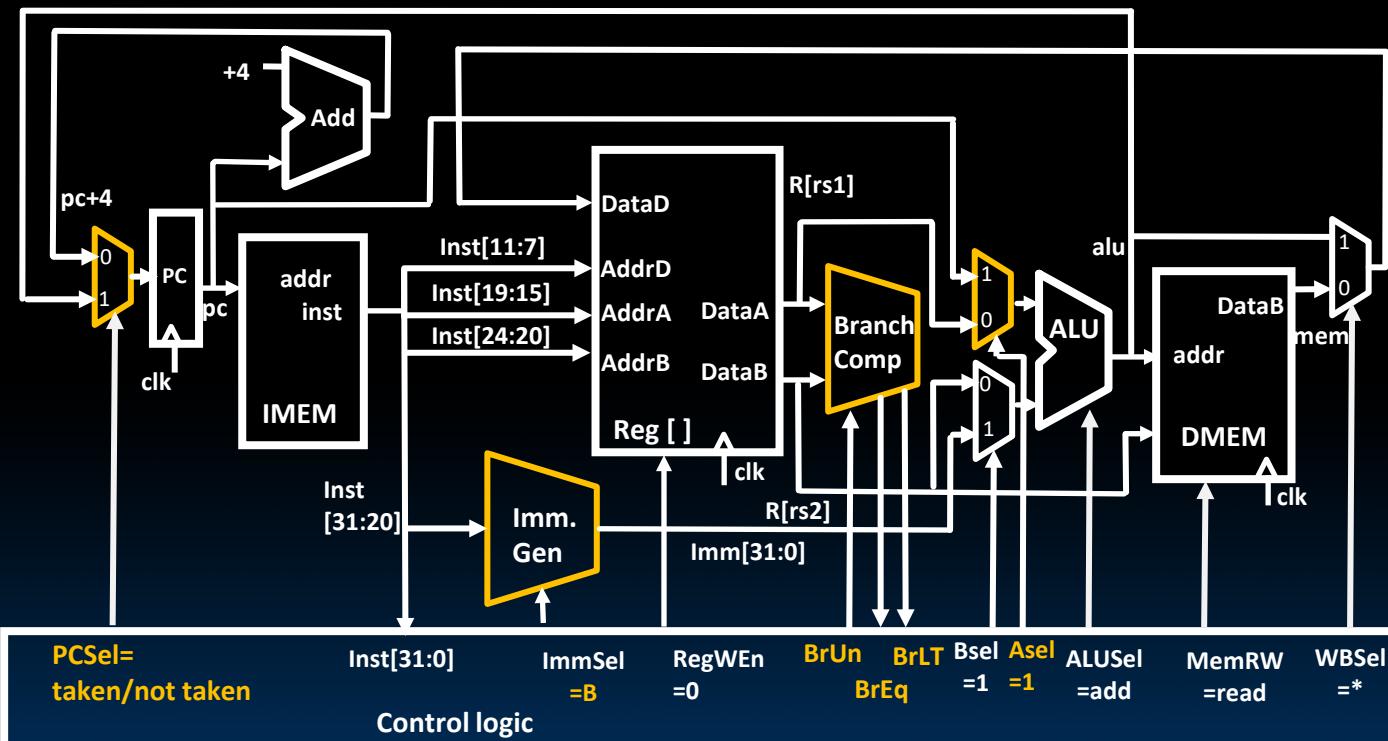
To Add Branches

- Different change to the state:

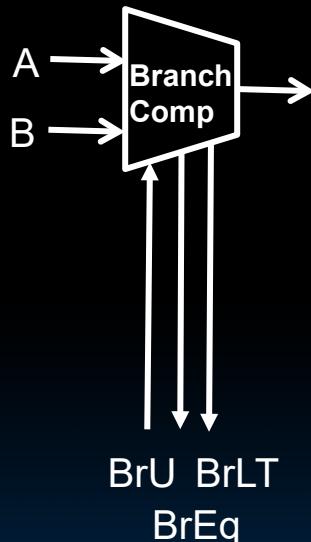
$$\text{PC} = \begin{cases} \text{PC} + 4, & \text{branch not taken} \\ \text{PC} + \text{immediate}, & \text{branch taken} \end{cases}$$

- Six branch instructions: **beq**, **bne**, **blt**, **bge**, **bltu**, **bgeu**
- Need to compute **PC + immediate** and to compare values of **rs1** and **rs2**
 - But have only one ALU – need more hardware

Adding Branches



Branch Comparator



BrEq = 1, if $A=B$

BrLT = 1, if $A < B$

BrUn = 1 selects unsigned comparison for **BrLT**,
0=signed

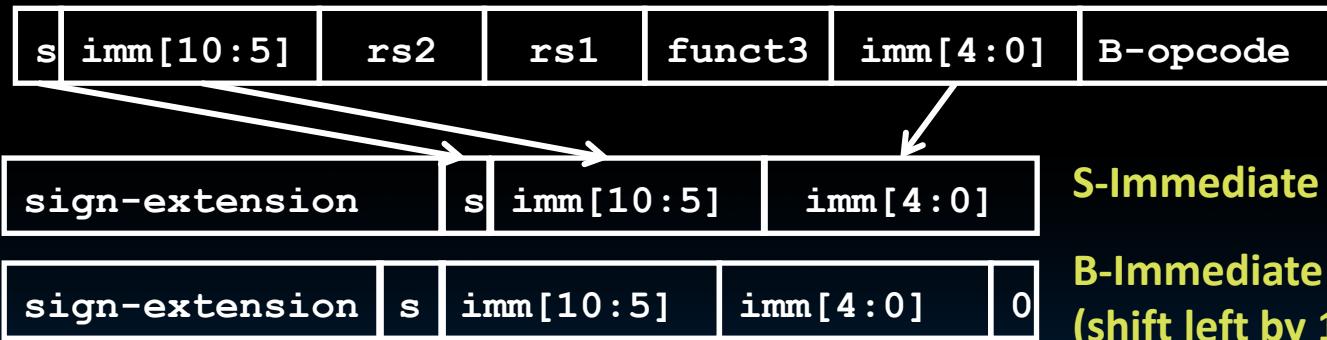
BGE branch: $A \geq B$, if $A < B$

$$\overline{A < B} = !(A < B)$$

Branch Immediates (In Other ISAs)

12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes

Standard approach: Treat immediate as in range -2048..+2047, then shift left by 1 bit to multiply by 2 for branches



Each instruction immediate bit can appear in one of two places in output immediate value – so need one 2-way mux per bit

Branch Immediates (In Other ISAs)

12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes

RISC-V approach: keep 11 immediate bits in fixed position in output value



Only one bit changes position between S and B, so only need a single-bit 2-way mux

RISC-V Immediate Encoding

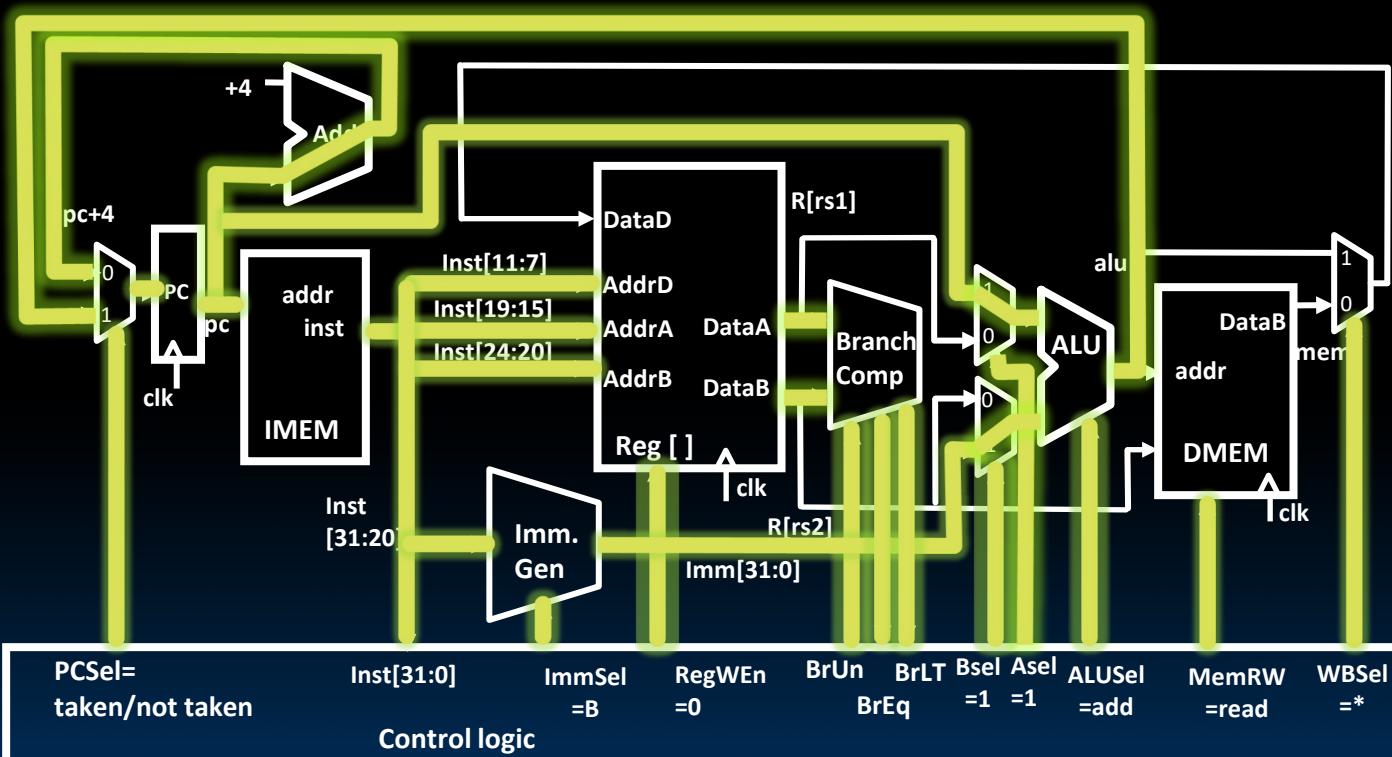
Instruction encodings, $\text{inst}[31:0]$											
31	30	25 24	20 19	15 14	12 11	8	7	6	0		
		imm[11:0]		rs1	funct3		rd		opcode		I-type
		imm[11:5]	rs2	rs1	funct3	imm[4:0]		opcode			S-type
		imm[12 10:5]	rs2	rs1	funct3	imm[4:1 11]		opcode			B-type

32-bit immediates produced, $\text{imm}[31:0]$											
31	25 24	12	11	10	5	4	1	0			
	-inst[31]-			inst[30:25]	inst[24:21]	inst[20]					I-imm.
	-inst[31]-			inst[30:25]	inst[11:8]	inst[7]					S-imm.
	-inst[31]-	inst[7]	inst[30:25]	inst[11:8]			0				B-imm.

Upper bits sign-extended from `inst[31]` always

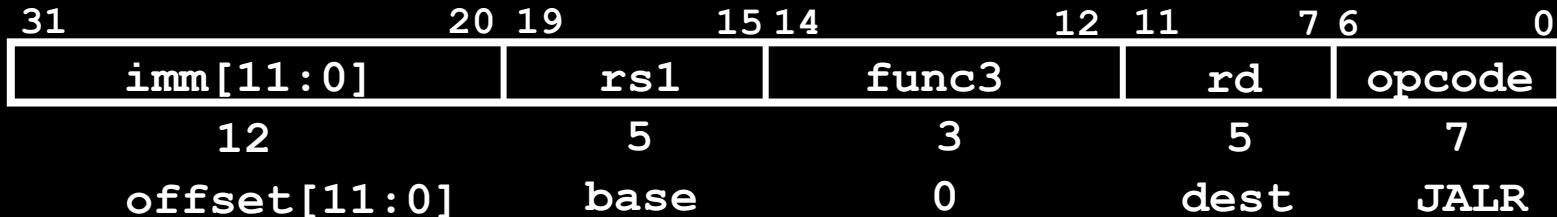
Only bit 7 of instruction changes role in immediate between S and B

Lighting Up Branch Path



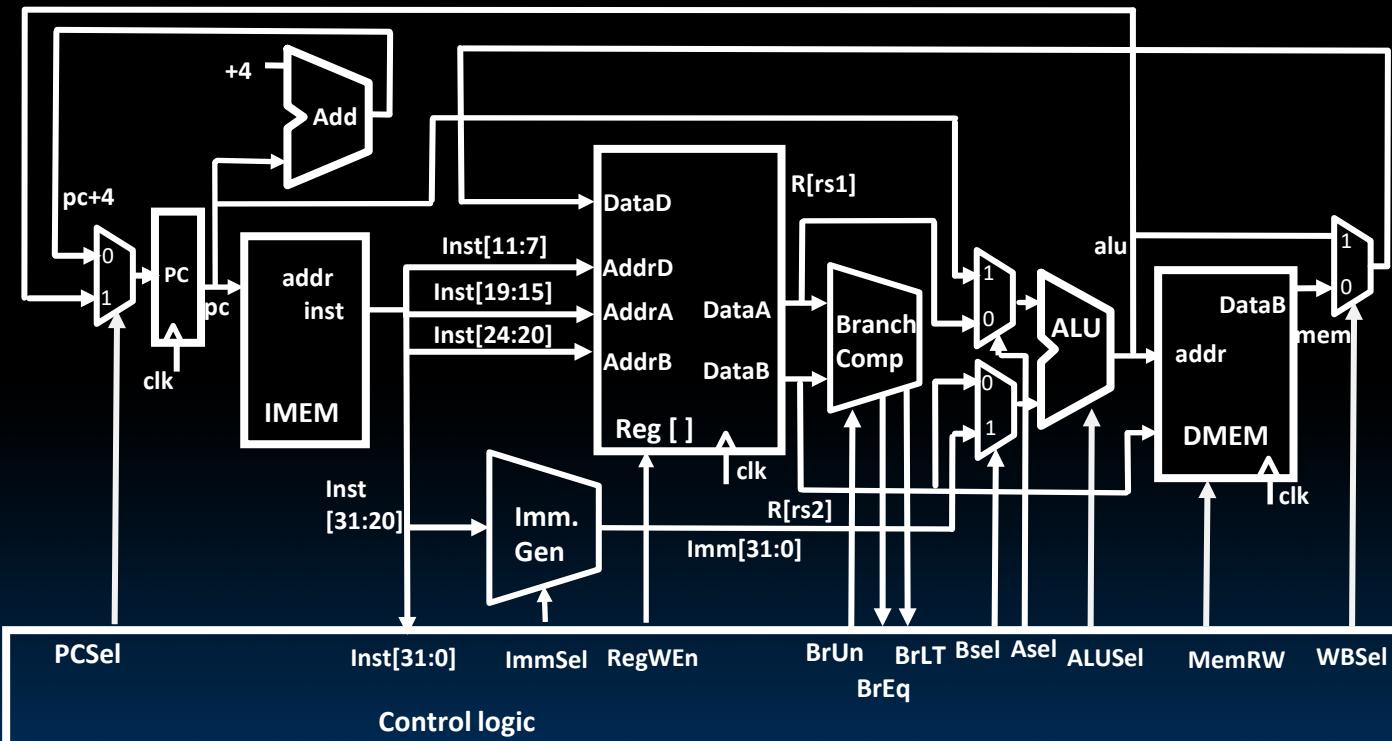
Adding JALR to Datapath

Let's Add JALR (I-Format)

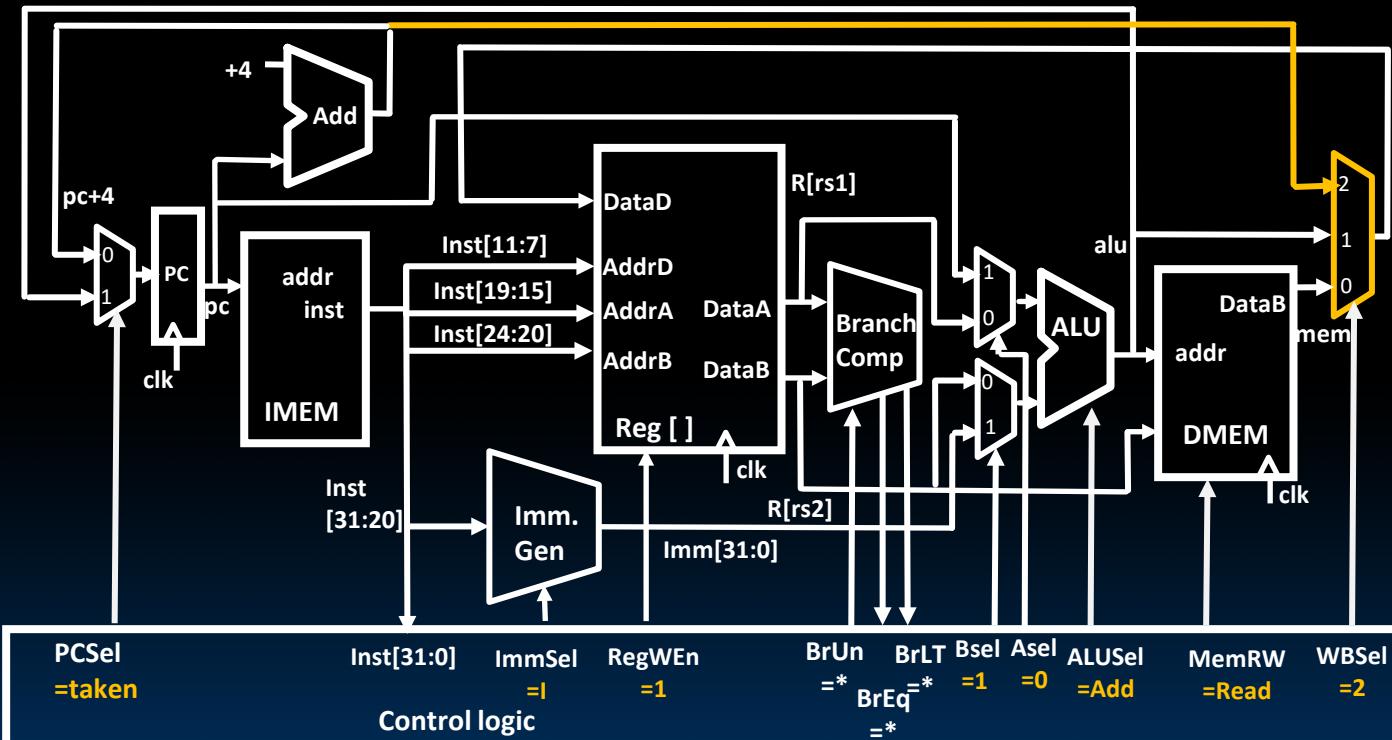


- **JALR rd, rs, immediate**
- Two changes to the state
 - Writes PC+4 to **rd** (return address)
 - Sets $PC = rs1 + \text{immediate}$
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - LSB is ignored

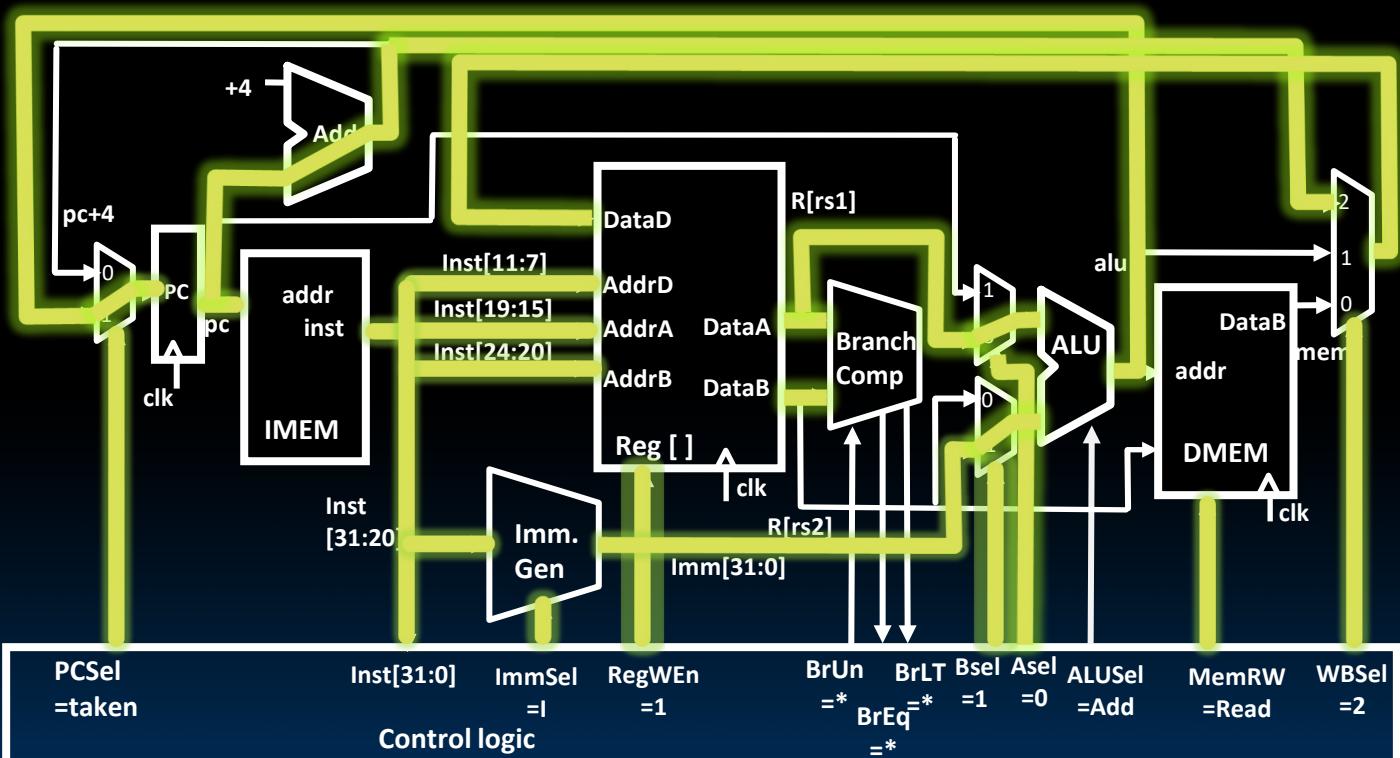
Datapath So Far, With Branches



Datapath With JALR

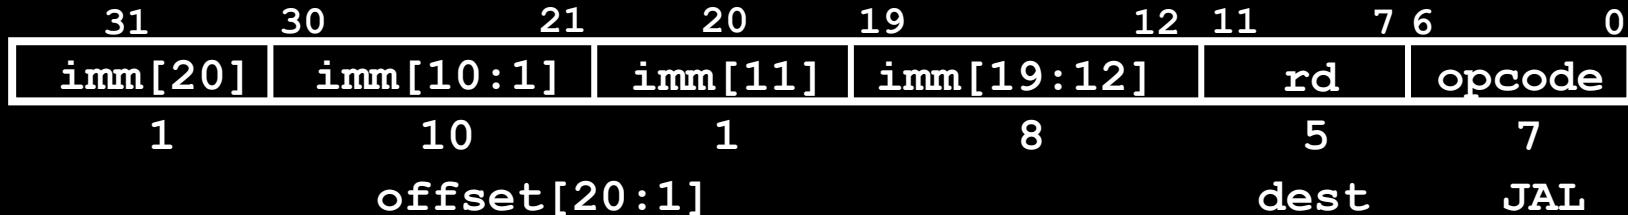


Datapath With JALR



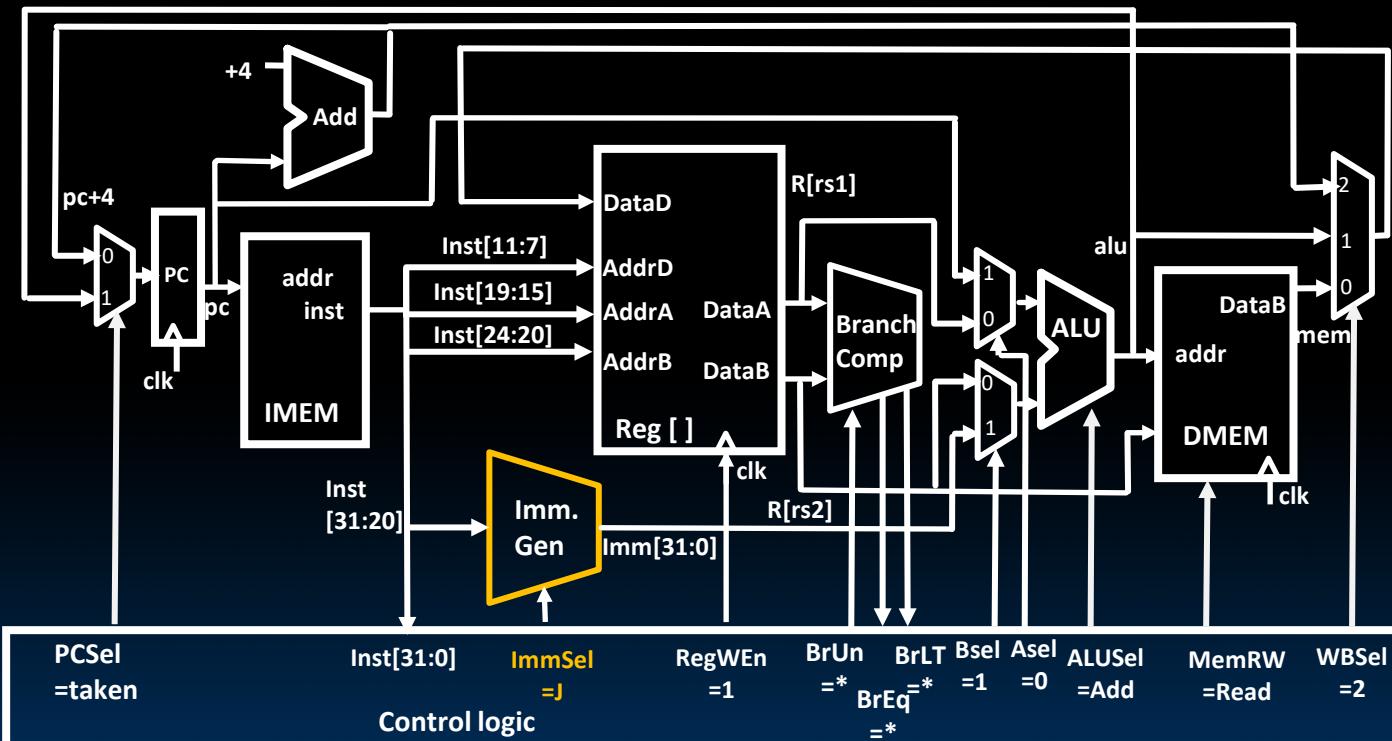
Adding JAL

J-Format for Jump Instructions

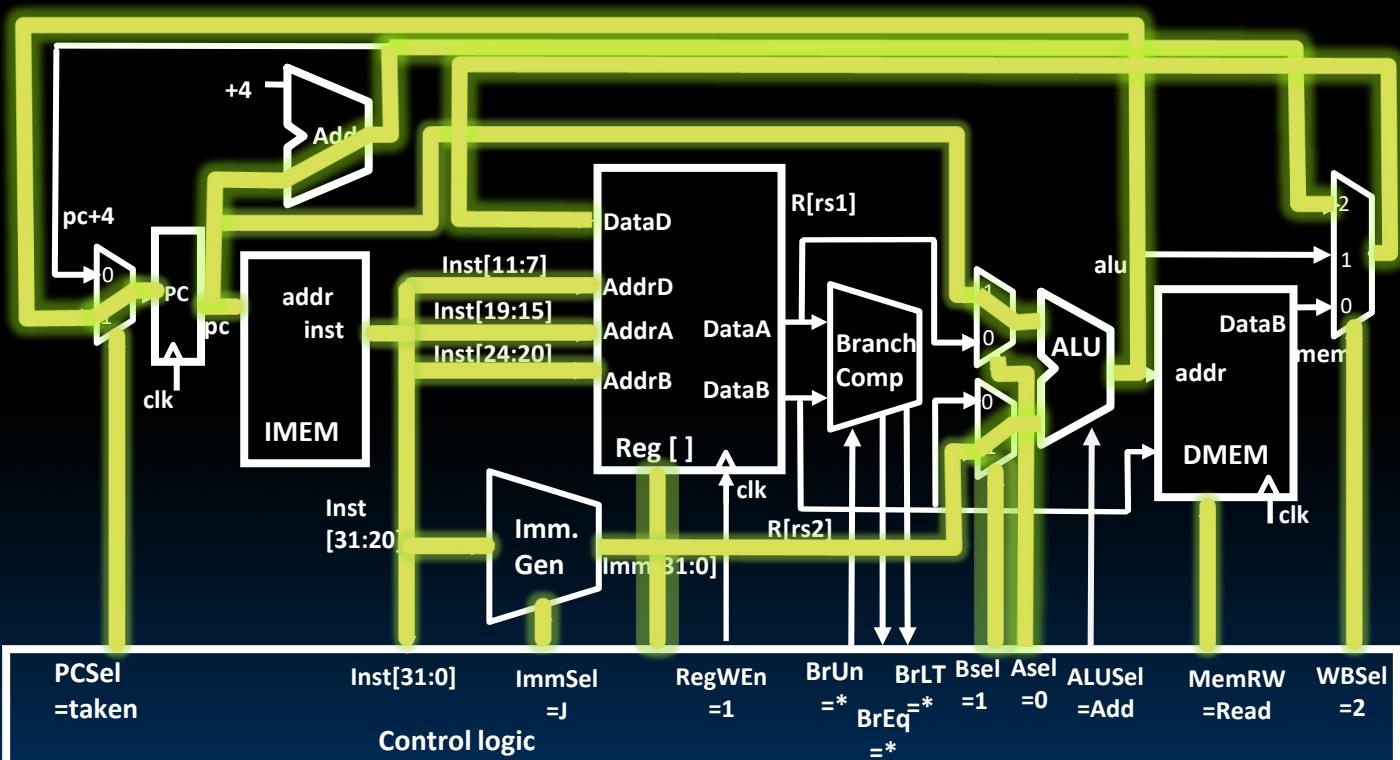


- Two changes to the state
 - `jal` saves PC+4 in register `rd` (the return address)
 - Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

Datapath with JAL

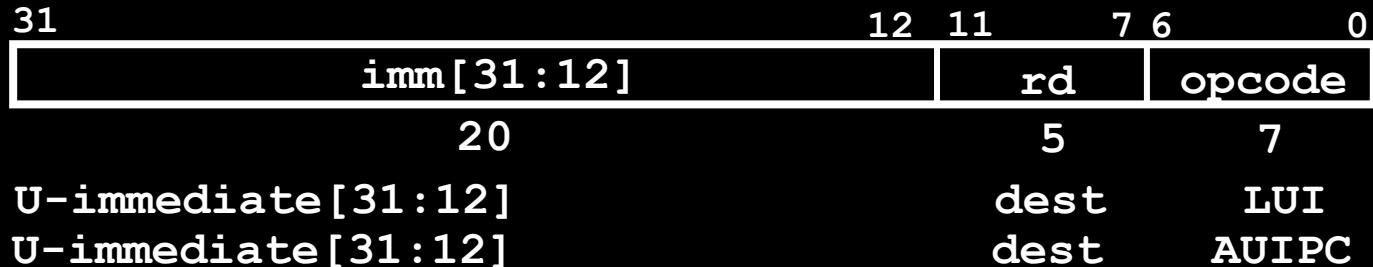


Light Up JAL Path



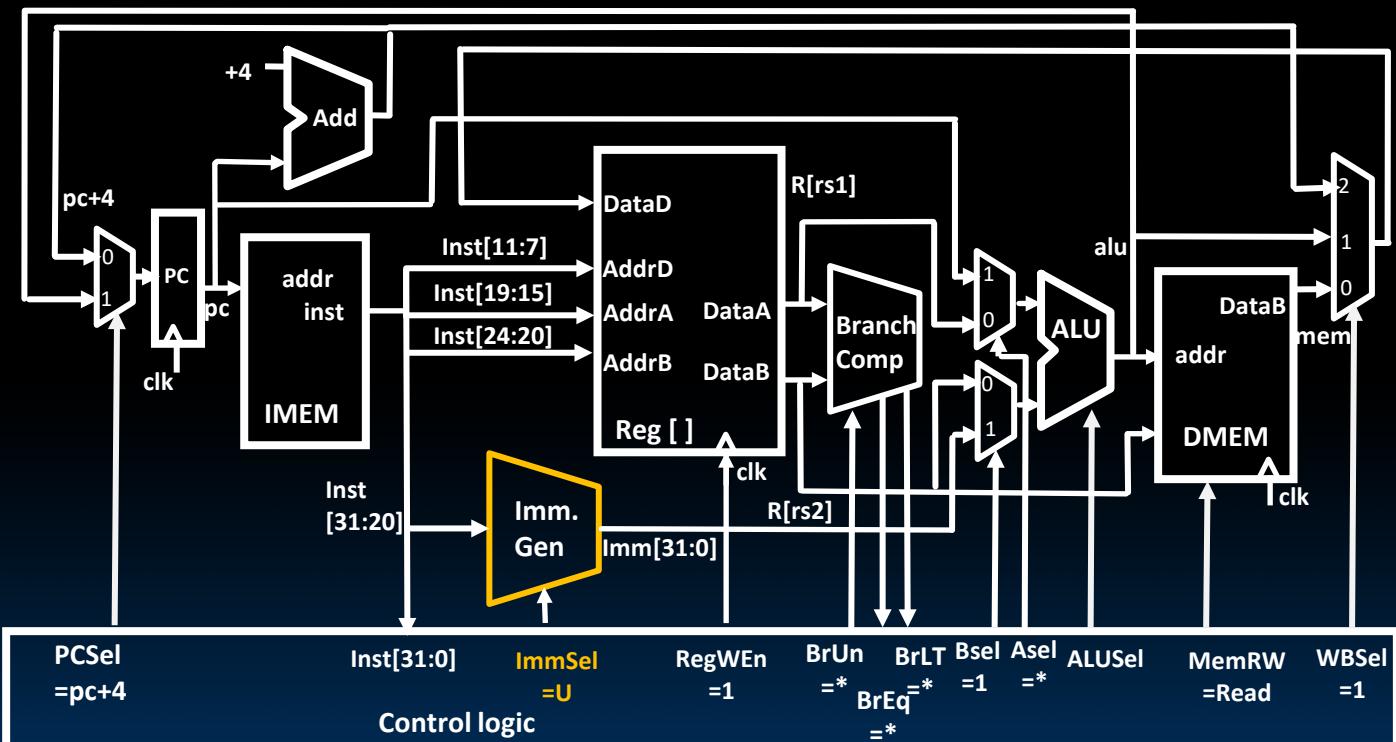
Adding U-Types

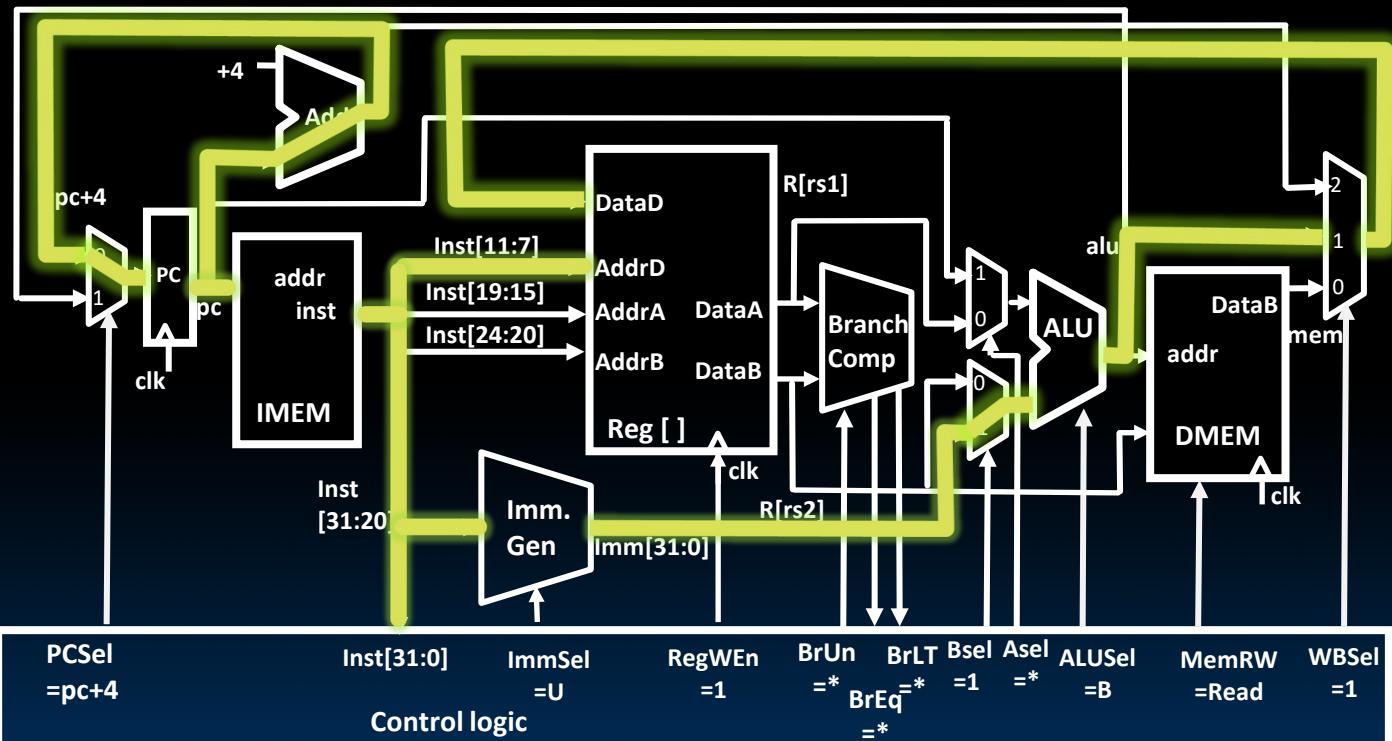
U-Format for “Upper Immediate” Instructions



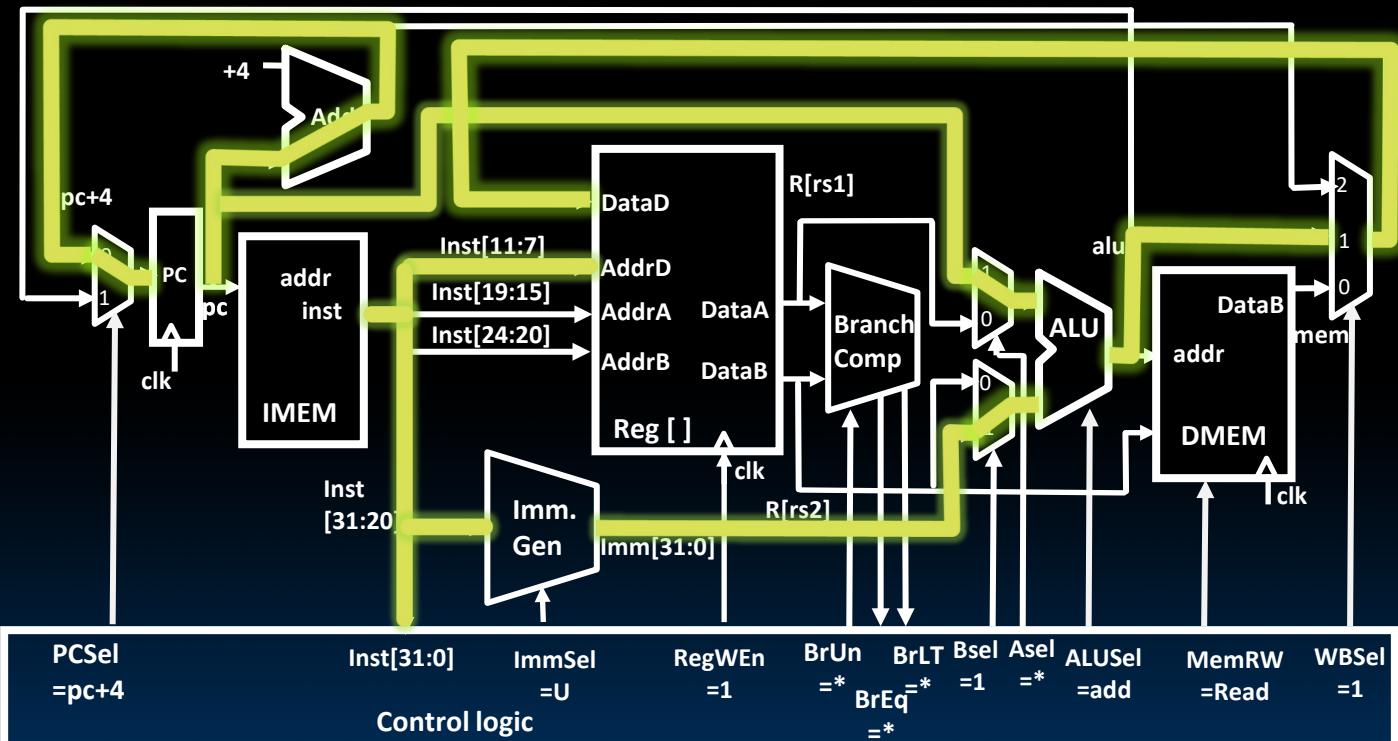
- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, **rd**
- Used for two instructions
 - **lui** – Load Upper Immediate
 - **auipc** – Add Upper Immediate to PC

Datapath With LUI, AUIPC



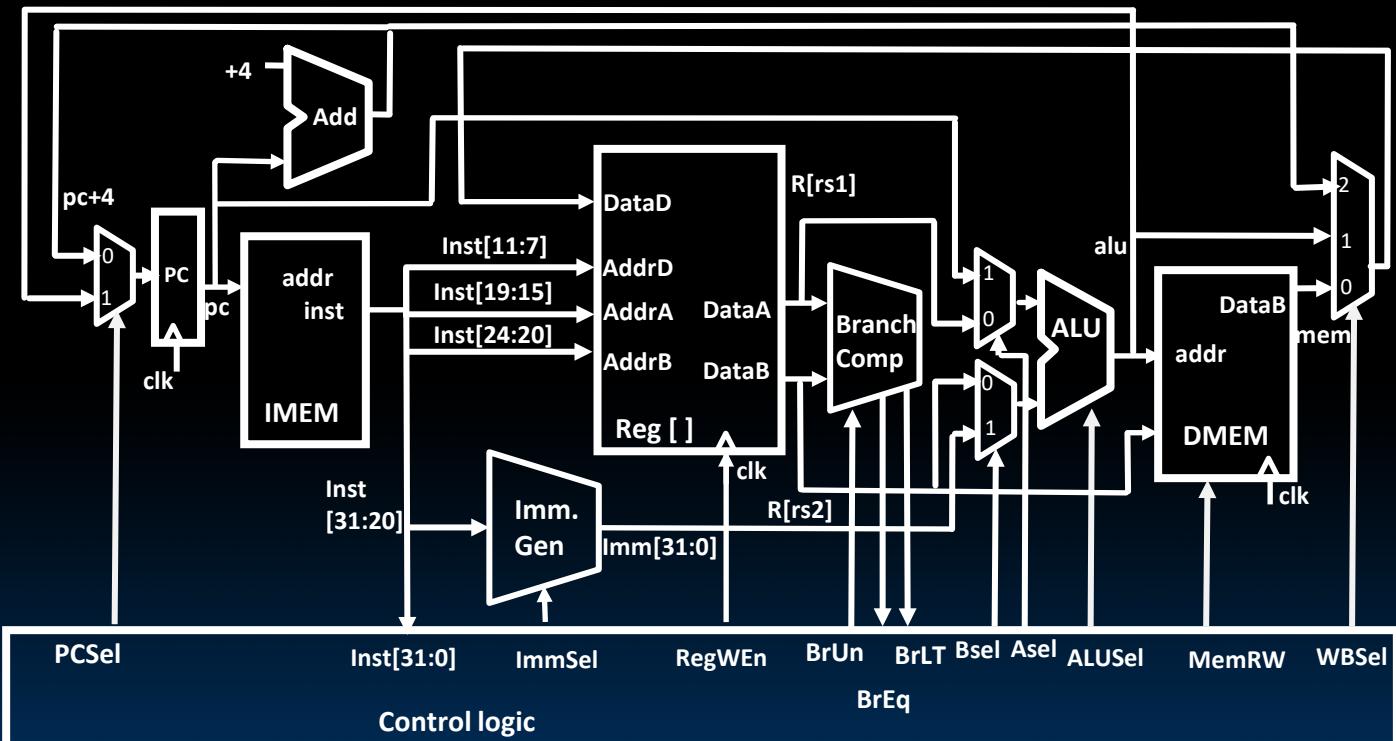


Lighting Up AUIPC



**“And In
Conclusion...”**

Complete RV32I Datapath!



Complete RV32I ISA!

OpenRISC RISC-V Reference Card

Base Integer Instructions: RV32I									
Category	Name	Fmt	RV32I Base		Category	Name	Fmt	RV32I Base	
Shifts	Shift Left Logical	R	SLL	rd,rs1,rs2	Loads	Load Byte	I	LB	rd,rs1,imm
	Shift Left Log. Imm.	I	SLLI	rd,rs1,shamt		Load Halfword	I	LH	rd,rs1,imm
	Shift Right Logical	R	SRL	rd,rs1,rs2		Load Byte Unsigned	I	LBU	rd,rs1,imm
	Shift Right Log. Imm.	I	SRLI	rd,rs1,shamt		Load Half Unsigned	I	LHU	rd,rs1,imm
	Shift Right Arithmetic	R	SRA	rd,rs1,rs2		Load Word	I	LW	rd,rs1,imm
	Shift Right Arith. Imm.	I	SRAI	rd,rs1,shamt	Stores	Store Byte	S	SB	rs1,rs2,imm
Arithmetic	ADD	R	ADD	rd,rs1,rs2		Store Halfword	S	SH	rs1,rs2,imm
	ADD Immediate	I	ADDI	rd,rs1,imm		Store Word	S	SW	rs1,rs2,imm
	SUBtract	R	SUB	rd,rs1,rs2	Branches	Branch =	B	BEQ	rs1,rs2,imm
	Load Upper Imm	U	LUI	rd,imm		Branch ≠	B	BNE	rs1,rs2,imm
Logical	Add Upper Imm to PC	U	AUIPC	rd,imm		Branch <	B	BLT	rs1,rs2,imm
	XOR	R	XOR	rd,rs1,rs2		Branch ≥	B	BGE	rs1,rs2,imm
	XOR Immediate	I	XORI	rd,rs1,imm		Branch < Unsigned	B	BLTU	rs1,rs2,imm
	OR	R	OR	rd,rs1,rs2		Branch ≥ Unsigned	B	BGEU	rs1,rs2,imm
	OR Immediate	I	ORI	rd,rs1,imm	Jump & Link	J&L	J	JAL	rd,imm
	AND	R	AND	rd,rs1,rs2		Jump & Link Register	I	JALR	rd,rs1,imm
	AND Immediate	I	ANDI	rd,rs1,imm					
Compare	Set <	R	SLT	rd,rs1,rs2	Synch	Synch thread	I	FENCE	
	Set < Immediate	I	SLTI	rd,rs1,imm					
	Set < Unsigned	R	SLTU	rd,rs1,rs2	Environment	CALL	I	ECALL	Not in 61C
	Set < Imm Unsigned	I	SLTIU	rd,rs1,imm		BREAK	I	EBREAK	

- We have designed a complete datapath
 - Capable of executing all RISC-V instructions in one cycle each
 - Not all units (hardware) used by all instructions
- 5 Phases of execution
 - IF, ID, EX, MEM, WB
 - Not all instructions are active in all phases
- Controller specifies how to execute instructions
 - We still need to design it